

Direct Rendering of Feature-based Skinning Deformations

PhD Thesis

Andreas-Alexandros Vasilakis

abasilak@cs.uoi.gr

Department of Computer Science & Engineering,
University of Ioannina, Greece

January 2014

Outline

- 1 Background and Preliminaries
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 S-buffer: Sparsity-aware multi-fragment rendering
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 k^+ -buffer: Fragment synchronized k -buffer
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Outline

- 1 **Background and Preliminaries**
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 S-buffer: Sparsity-aware multi-fragment rendering
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 k^+ -buffer: Fragment synchronized k -buffer
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Problem statement

Direct rendering skinned approximations of self-crossing deformable objects

- ① [skeletal animation] (mainly driven by segmentation)
- ② [interactive rasterization]
 - multi-fragment rendering
 - self-intersection surfaces

Problem statement

Direct rendering skinned approximations of self-crossing deformable objects

- ① [skeletal animation] (mainly driven by segmentation)
- ② [interactive rasterization]
 - multi-fragment rendering
 - self-intersection surfaces

Wide range of applications

- animation compression (skinning), skeleton extraction, deformation transfer, volume preservation.
- post-editing, lod selection, collision detection, game engines.
- multi-fragment rendering effects:
 - interior visualization (transparency, clipping, capping, volume rendering)
 - self-trimming and constructive solid geometry (CSG) operations
 - and more: shadows, hair rendering, global illumination, voxelization...

Graphics rendering pipeline

[Main function]: generate (*render*) a 2D image given a virtual camera, 3D objects, light sources, textures, etc...

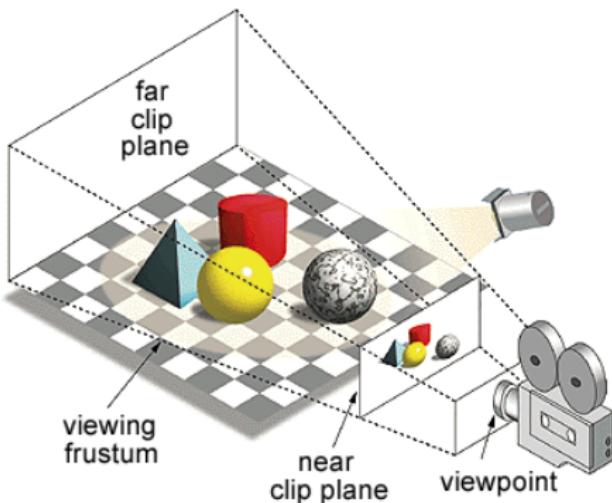


Figure: The rendering process of using this pipeline.

Graphics rendering pipeline

Division into three conceptual stages:

- ➊ [Application]: fed next stage with the *rendering primitives*

Graphics rendering pipeline

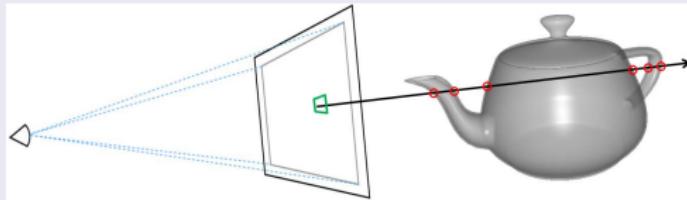
Division into three conceptual stages:

- ① [Application]: fed next stage with the *rendering primitives*
- ② [Geometry]: per-vertex & per-polygon operations

Graphics rendering pipeline

Division into three conceptual stages:

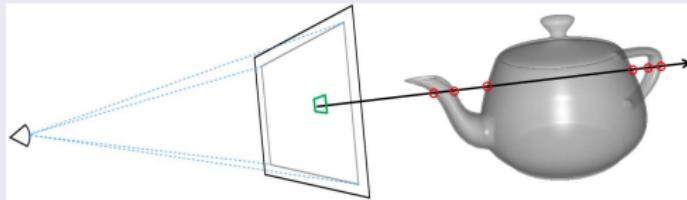
- ① [Application]: fed next stage with the *rendering primitives*
- ② [Geometry]: per-vertex & per-polygon operations
- ③ [Rasterizer]: generates the final image by computing the colors for the pixels covered by visible objects.



Graphics rendering pipeline

Division into three conceptual stages:

- ① [Application]: fed next stage with the *rendering primitives*
- ② [Geometry]: per-vertex & per-polygon operations
- ③ [Rasterizer]: generates the final image by computing the colors for the pixels covered by visible objects.

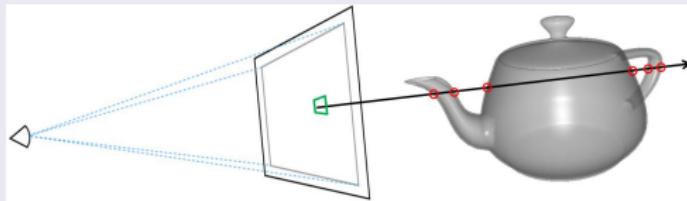


- ① [Scan-conversion]: fragment generation (vertex data interpolation)

Graphics rendering pipeline

Division into three conceptual stages:

- ① [Application]: fed next stage with the *rendering primitives*
- ② [Geometry]: per-vertex & per-polygon operations
- ③ [Rasterizer]: generates the final image by computing the colors for the pixels covered by visible objects.

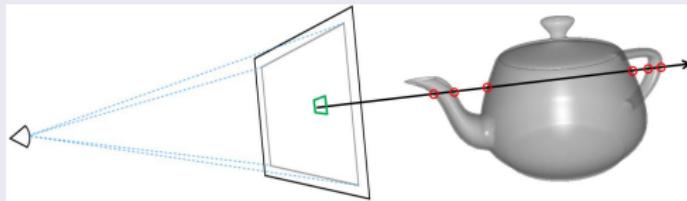


- ① [Scan-conversion]: fragment generation (vertex data interpolation)
- ② [Shading]: stores the color(s) and depth for each pixel in buffers

Graphics rendering pipeline

Division into three conceptual stages:

- ① [Application]: fed next stage with the *rendering primitives*
- ② [Geometry]: per-vertex & per-polygon operations
- ③ [Rasterizer]: generates the final image by computing the colors for the pixels covered by visible objects.



- ① [Scan-conversion]: fragment generation (vertex data interpolation)
- ② [Shading]: stores the color(s) and depth for each pixel in buffers
- ③ [Merging]: resolves visibility via Z-testing and performs optionally various (*blending, stencil test*) operations.

Background

Pose-to-pose segmentation

Pose-to-pose skinning

S-buffer

Eliminating z-fighting flaws

k^+ -buffer

Rendering self-trimmed solids

Conclusions & Future Work

Problem Statement

Graphics Rendering Pipeline

Framework Overview

Mesh Animation (MA)

Segmenting MA

Skinning MA

Self-intersected MA

Multi-fragment Rendering

Framework overview

System components

Framework overview

System components

{*application-stage*}

- ① [Editing]: deforms an arbitrary static 3D object - mesh animation (MA)

Framework overview

System components

{*application-stage*}

- ① [Editing]: deforms an arbitrary static 3D object - mesh animation (MA)

{*geometry-stage*}

- ② [Segmentation]: partitions MA into parts with similar motion features

Framework overview

System components

- { *application-stage* }
 - ① [Editing]: deforms an arbitrary static 3D object - mesh animation (MA)
- { *geometry-stage* }
 - ② [Segmentation]: partitions MA into parts with similar motion features
 - ③ [Skinning]: transforms MA into a compact skeletal representation

Framework overview

System components

- { *application-stage* }
 - ① [Editing]: deforms an arbitrary static 3D object - mesh animation (MA)
- { *geometry-stage* }
 - ② [Segmentation]: partitions MA into parts with similar motion features
 - ③ [Skinning]: transforms MA into a compact skeletal representation
- { *rasterizer-stage* }
 - ④ [Fragment storage]: capture generated fragments in correct depth order

Framework overview

System components

- { *application-stage* }
 - ① [Editing]: deforms an arbitrary static 3D object - mesh animation (MA)
- { *geometry-stage* }
 - ② [Segmentation]: partitions MA into parts with similar motion features
 - ③ [Skinning]: transforms MA into a compact skeletal representation
- { *rasterizer-stage* }
 - ④ [Fragment storage]: capture generated fragments in correct depth order
 - ⑤ [Self-trimming]: generates the trimmed surface (watertight solid)

Framework overview

System components

- { *application-stage* }
 - ① [Editing]: deforms an arbitrary static 3D object - mesh animation (MA)
- { *geometry-stage* }
 - ② [Segmentation]: partitions MA into parts with similar motion features
 - ③ [Skinning]: transforms MA into a compact skeletal representation
- { *rasterizer-stage* }
 - ④ [Fragment storage]: capture generated fragments in correct depth order
 - ⑤ [Self-trimming]: generates the trimmed surface (watertight solid)
 - ⑥ [Extra effects]: performs interior visualization or CSG operations

Mesh animation (MA)

A sequence of key-frames $\{p_t, t = 1, \dots, k\}$, representing how a static 3D shape p_0 is evolving through time (may also self-intersect).

Mesh animation (MA)

A sequence of key-frames $\{p_t, t = 1, \dots, k\}$, representing how a static 3D shape p_0 is evolving through time (may also self-intersect).

Figure: Horse, skirt and self-crossing sphere mesh animations.

Mesh animation (MA)

A sequence of key-frames $\{p_t, t = 1, \dots, k\}$, representing how a static 3D shape p_0 is evolving through time (may also self-intersect).

Production

- ① [automatic]: scanning machinery, multiple video cameras
- ② [human intervention]: editing tools (skeleton, free-form)

Mesh animation (MA)

A sequence of key-frames $\{p_t, t = 1, \dots, k\}$, representing how a static 3D shape p_0 is evolving through time (may also self-intersect).

Production

- ① [automatic]: scanning machinery, multiple video cameras
- ② [human intervention]: editing tools (skeleton, free-form)

Categories

- ① time-varying versus [temporal-coherent]
- ② [off-line] vs [real-time] (streamed or dynamically generated)

Mesh animation (MA)

A sequence of key-frames $\{p_t, t = 1, \dots, k\}$, representing how a static 3D shape p_0 is evolving through time (may also self-intersect).

Production

- ① [automatic]: scanning machinery, multiple video cameras
- ② [human intervention]: editing tools (skeleton, free-form)

Categories

- ① time-varying versus [temporal-coherent]
- ② [off-line] vs [real-time] (streamed or dynamically generated)

Types

- ① [rigid] (human, quadrupeds)
 - ② [highly-deformable] (cloth, facial)
- } [hybrid] animation

Segmenting animated meshes

[Main goal]: partitions the animated mesh into clusters $\{C^0, \dots, C^I\}$ with similar motion features

Segmenting animated meshes

[Main goal]: partitions the animated mesh into clusters $\{C^0, \dots, C^I\}$ with similar motion features

Figure: Segmenting a horse mesh animation into 42 parts.

Segmenting animated meshes

[Main goal]: partitions the animated mesh into clusters $\{C^0, \dots, C^I\}$ with similar motion features

Global segmentation methods

- ① [clustering technique]: k-means, mean shift, hierarchical, spectral, ...
- ② [feature space]: euclidean, geodesic & angular distances, ...

Segmenting animated meshes

[Main goal]: partitions the animated mesh into clusters $\{C^0, \dots, C^I\}$ with similar motion features

Global segmentation methods

- ➊ [clustering technique]: k-means, mean shift, hierarchical, spectral, ...
- ➋ [feature space]: euclidean, geodesic & angular distances, ...

Limitations

- focus on detecting segments with mostly *rigid* behavior
- fail when feature space is *flat* or *anisotropic*
- work only with *off-line* animations (not-editable)
- cannot offer smooth multi-resolution & variable¹ segmentations

¹ [Arcila et al, GM'13]

Skeletal animation

Use a simpler structure ([skeleton](#)) to animate mesh instead of editing each primitive independently.

Skeletal animation

Use a simpler structure (**skeleton**) to animate mesh instead of editing each primitive independently.

Representation

- (a) reference 3D surface mesh
- (b) endoskeleton (joints, bones) attached to the skin (*rigging*)
- (c) anatomy layer (musculature)
- (d) direct manipulation, inverse kinematics or motion data



Mesh Skinning

- efficient (GPU-accelerated), widely used from game engines

Mesh Skinning

- efficient (GPU-accelerated), widely used from game engines

Skin deformation algorithm

- reference 3D surface mesh (rest-pose)
- a list of joint transformations (no actual skeleton needed)
- two lists of vertex indices and influencing weights

Mesh Skinning

- efficient (GPU-accelerated), widely used from game engines

Skin deformation algorithm

- reference 3D surface mesh (rest-pose)
- a list of joint transformations (no actual skeleton needed)
- two lists of vertex indices and influencing weights

Predominant skinning techniques

- Linear blend skinning (**LBS¹**): fastest, elongations & interpenetrations
- Dual quaternion skinning (**DQS²**): rigid motion, less storage, extra cost

¹ [Magnenat-Thalmann et al, GI'88], ² [Kavan et al, TOG'08]

Mesh Skinning

- efficient (GPU-accelerated), widely used from game engines

Skin deformation algorithm

- reference 3D surface mesh (rest-pose)
- a list of joint transformations (no actual skeleton needed)
- two lists of vertex indices and influencing weights

Predominant skinning techniques

- Linear blend skinning (**LBS¹**): fastest, elongations & interpenetrations
- Dual quaternion skinning (**DQS²**): rigid motion, less storage, extra cost

Approximate MA with skinning

- **[Goal]:** animation compression (data reduction)
- **[Applications]:** GPU, collision detection, lod selection, pose editing

Self-intersected Mesh Animations

Geometry processing pipeline:

- [deformation], smoothing, subdivision and decimation
- [skinning approximation]

Self-intersected Mesh Animations

Geometry processing pipeline:

- [deformation], smoothing, subdivision and decimation
- [skinning approximation]

Problems

- ① Given a *self-crossing surface* (SCS) → [trim]
- ② Given a SCS derived by deforming an initial non-SCS surface → [trim]



Self-intersected Mesh Animations

Geometry processing pipeline:

- [deformation], smoothing, subdivision and decimation
- [skinning approximation]

Problems

- ① Given a *self-crossing surface* (SCS) → [trim]
- ② Given a SCS derived by deforming an initial non-SCS surface → [trim]



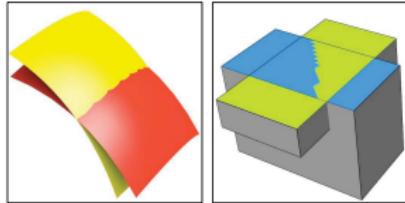
[geometry space] versus [image space]

Multi-fragment rendering

Z-fighting phenomenon

more

- two or more primitives have same depth
- intersecting or overlapping surfaces

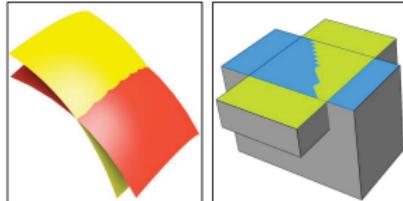


Multi-fragment rendering

Z-fighting phenomenon

more

- two or more primitives have same depth
- intersecting or overlapping surfaces



Storing multiple fragments

- multi-pass depth peeling (F2B¹, Dual², Bucket³)

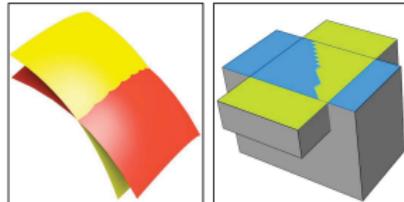
¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, HPG'09]

Multi-fragment rendering

Z-fighting phenomenon

more

- two or more primitives have same depth
- intersecting or overlapping surfaces



Storing multiple fragments

- multi-pass depth peeling (F2B¹, Dual², Bucket³) [z-fighting flaws]

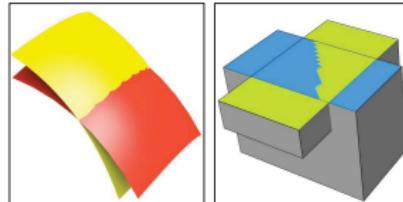
¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, HPG'09]

Multi-fragment rendering

Z-fighting phenomenon

more

- two or more primitives have same depth
- intersecting or overlapping surfaces



Storing multiple fragments

- multi-pass depth peeling ([F2B¹](#), [Dual²](#), [Bucket³](#)) [z-fighting flaws]
- hardware buffers ([A-buffer⁴](#), [k-buffer⁵](#))

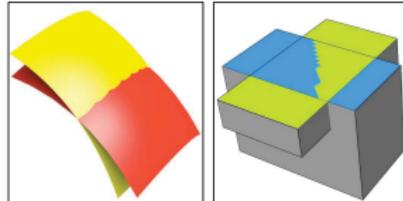
⁴ [Carpenter, SIGGRAPH'84], ⁵ [Bavoil et al, I3D'07]

Multi-fragment rendering

Z-fighting phenomenon

more

- two or more primitives have same depth
- intersecting or overlapping surfaces



Storing multiple fragments

- multi-pass depth peeling ([F2B¹](#), [Dual²](#), [Bucket³](#)) [z-fighting flaws]
- hardware buffers ([A-buffer⁴](#), [k-buffer⁵](#)) [fragment congestion & overflow]

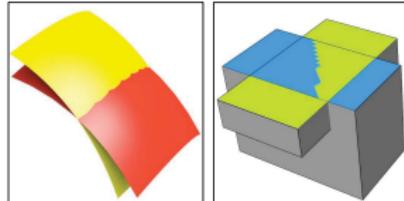
⁴ [Carpenter, SIGGRAPH'84], ⁵ [Bavoil et al, I3D'07]

Multi-fragment rendering

Z-fighting phenomenon

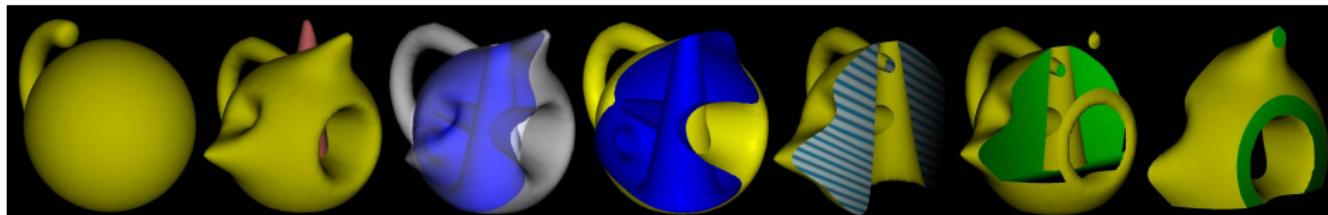
more

- two or more primitives have same depth
- intersecting or overlapping surfaces



Storing multiple fragments

- multi-pass depth peeling ([F2B¹](#), [Dual²](#), [Bucket³](#)) [z-fighting flaws]
- hardware buffers ([A-buffer⁴](#), [k-buffer⁵](#)) [fragment congestion & overflow]



Outline

- 1 Background and Preliminaries
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 S-buffer: Sparsity-aware multi-fragment rendering
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 k^+ -buffer: Fragment synchronized k -buffer
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Revising segmentation of MAs

Segmentation Problem

[go back](#)

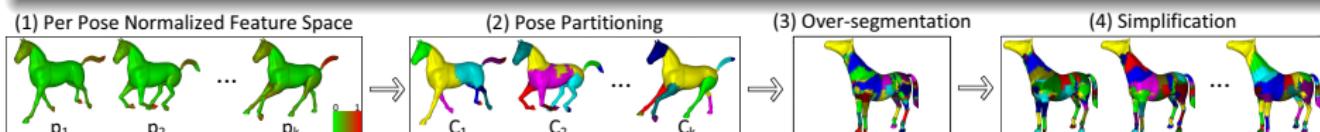
- handles all types and categories of MAs
- supports multi-resolution and variable segmentations
- efficiency and quality

Revising segmentation of MAs

go back

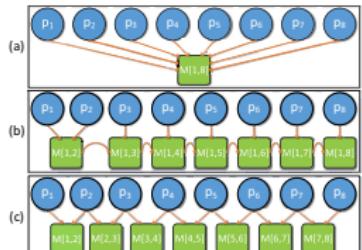
Segmentation Problem

- handles all types and categories of MAs
- supports multi-resolution and variable segmentations
- efficiency and quality



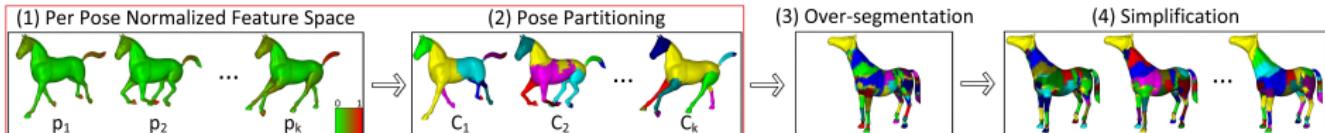
Our contribution†

- ① computes **per-pose partitioning** (feature)
- ② computes **over-segmentation** (OS)
- ③ performs progressive **simplification**

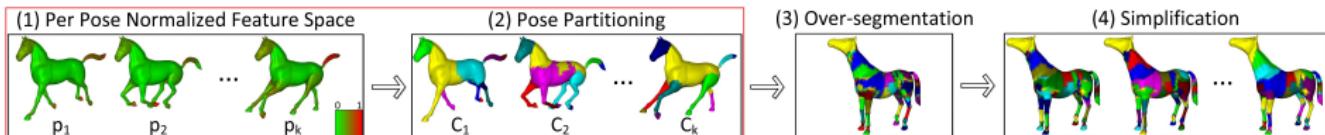


† A. Vasilakis & I. Fudos, *Pose Partitioning for Multi-resolution Segmentation of Arbitrary Mesh Animations*, Computer Graphics Forum (Proceedings of Eurographics 2014), vol. 33(2):?-?

(1,2) Per-pose segmentation



(1,2) Per-pose segmentation

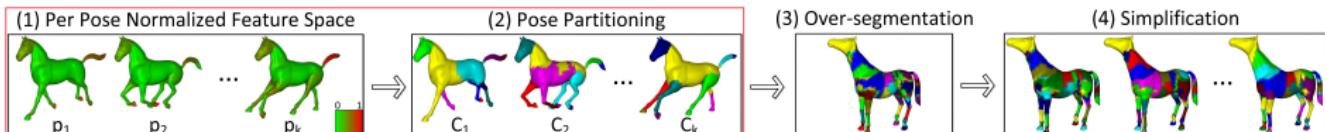


(1) Feature space computation

- a deformation measure from reference (p_0 or p_{t-1}) to target pose (p_t)
- deformation gradient \rightarrow [rotation angles]¹

¹ [Sumner et al., TOG'04] , ² [Gunther et al., CGF'06]

(1,2) Per-pose segmentation



(1) Feature space computation

- a deformation measure from reference (p_0 or p_{t-1}) to target pose (p_t)
- deformation gradient \rightarrow [rotation angles]¹

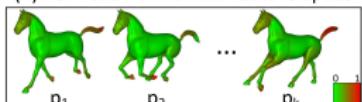
(2) Partitioning computation

- numerous clustering methods (in [parallel])
- [top-down hierarchical clustering]²
- [output]: $C_i = \{C_i^0, \dots, C_i^{l_i}\}, \forall p_i \in [1, k]$

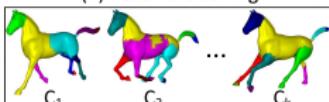
¹ [Sumner et al., TOG'04] , ² [Gunther et al., CGF'06]

(3) Pose partitioning-aware OS

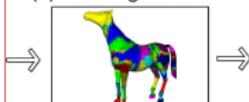
(1) Per Pose Normalized Feature Space



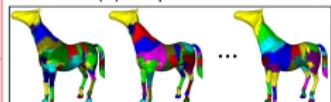
(2) Pose Partitioning



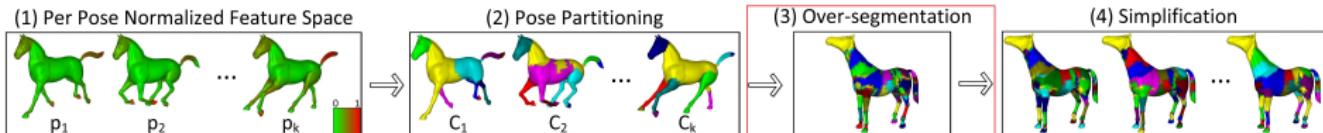
(3) Over-segmentation



(4) Simplification



(3) Pose partitioning-aware OS

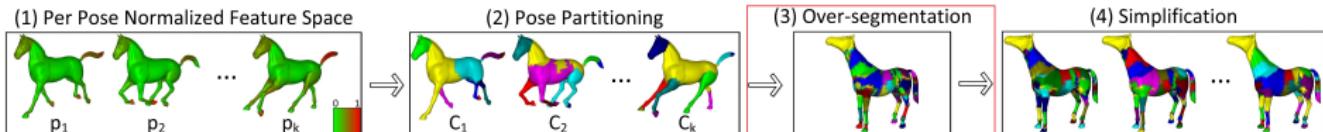


Clustering Animation Vector (CAV)

- $cav(v) = \{m_1, m_2, \dots, m_k\}, m_j \in 1, \dots, |C_j|, \forall v \in V$
- [Pruned-BFS]: detect connected vertices with same CAV - $O(|V|)$
- [output]: $OS(A) = \{S_0, S_1, \dots, S_I\}$

alg

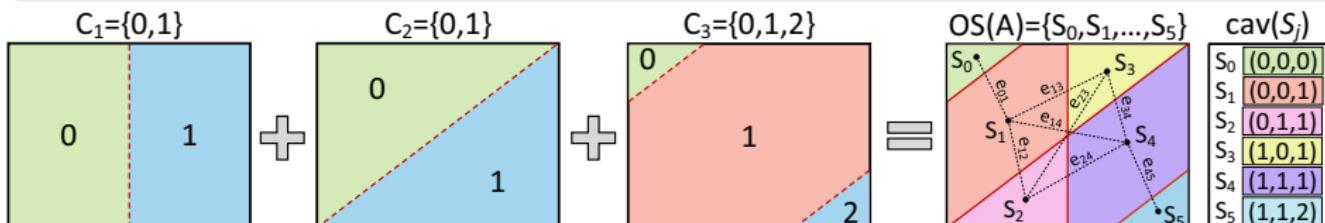
(3) Pose partitioning-aware OS



Clustering Animation Vector (CAV)

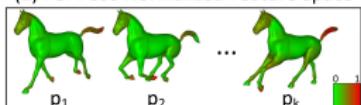
- $cav(v) = \{m_1, m_2, \dots, m_k\}, m_j \in 1, \dots, |C_j|, \forall v \in V$
- [Pruned-BFS]: detect connected vertices with same CAV - $O(|V|)$
- [output]: $OS(A) = \{S_0, S_1, \dots, S_l\}$

alg

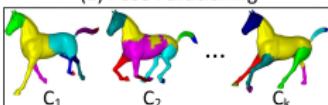


(4) Progressive decimation of OS

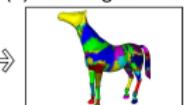
(1) Per Pose Normalized Feature Space



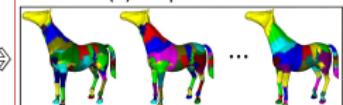
(2) Pose Partitioning



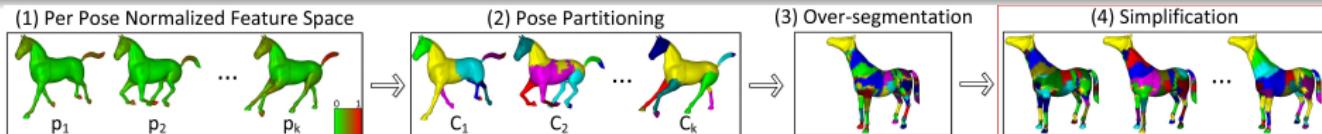
(3) Over-segmentation



(4) Simplification



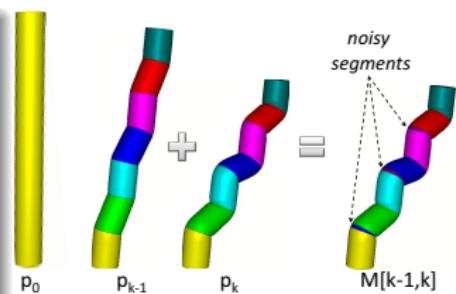
(4) Progressive decimation of OS



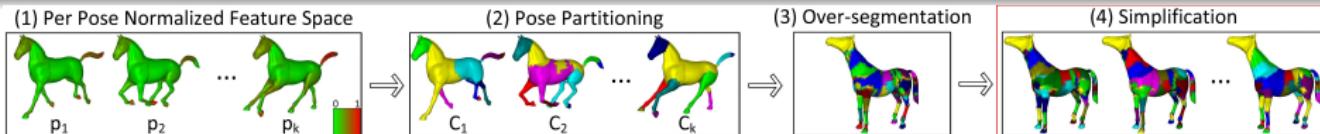
Pose-to-pose cleaning $R_i(h)$, $h \in [0, 1]$

alg

- $p_k \longmapsto, \dots, \longmapsto p_1$
- reduction rule for (S_A, S_B) :
 - $a(S_A) \leq h \cdot a(C_i^{cav(S_A)[i]})$
 - $a(S_A) \leq h \cdot a(\text{segment}(S_A, i - 1))$
 - $S_B \in N(S_A) \wedge cav(S_A)[1, i - 1] = cav(S_B)[1, i - 1] \wedge cav(S_A)[i] \neq cav(S_B)[i]$
 - $a(S_B) > a(S_A)$

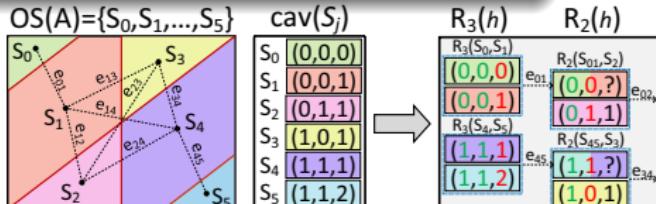
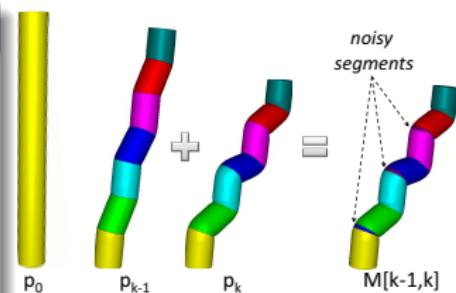


(4) Progressive decimation of OS



Pose-to-pose cleaning $R_i(h)$, $h \in [0, 1]$

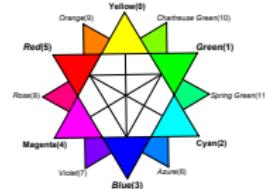
- $p_k \longmapsto, \dots, \longmapsto p_1$
- reduction rule for (S_A, S_B) :
 - $a(S_A) \leq h \cdot a(C_i^{cav(S_A)[i]})$
 - $a(S_A) \leq h \cdot a(\text{segment}(S_A, i - 1))$
 - $S_B \in N(S_A) \wedge cav(S_A)[1, i - 1] = cav(S_B)[1, i - 1] \wedge cav(S_A)[i] \neq cav(S_B)[i]$
 - $a(S_B) > a(S_A)$



Applications (1)

Smooth Visualization of Cluster Transitions

- minimize “close” colors assigned to neighbors
- propagate colors between consecutive frames



Applications (1)

Smooth Visualization of Cluster Transitions

- minimize “close” colors assigned to neighbors
- propagate colors between consecutive frames

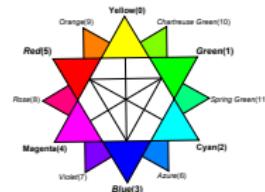


Figure: (left) initial painting and (right) color propagation.

Applications (2)

Multi-resolution segmentation

- High \rightarrow low: adjusting parameter $h : 0 \mapsto 1$
- Bounded by resolutions of $[C_1, OS]$ - [High similarity] ?
- Efficient: do not depend on mesh size

Applications (2)

Multi-resolution segmentation

- High \mapsto low: adjusting parameter $h : 0 \mapsto 1$
- Bounded by resolutions of $[C_1, OS]$ - [High similarity] ?
- Efficient: do not depend on mesh size

Applications (3)

back

Variable (temporal consistent) segmentation

- (a) remain faithful to past-data }
(b) alter to future data } [pair merging] $(C_{t-1}, C_t) \mapsto$ real-time

Applications (3)

Variable (temporal consistent) segmentation

back

- (a) remain faithful to past-data }
(b) alter to future data } [pair merging] $(C_{t-1}, C_t) \longmapsto$ real-time

Figure: (left) color propagation and (right) variable segmentation.

Applications (4)

back

Real-time segmentation

- streamed/dynamic MA
 - out-of-core algorithm
- $\left. \begin{array}{c} \\ \end{array} \right\}$ [merge] $(OS(0, t - 1), C_t) \longmapsto \text{cost of } C_t$

Applications (4)

back

Real-time segmentation

- streamed/dynamic MA
 - out-of-core algorithm
- $\left. \right\} \text{[merge]} (OS(0, t - 1), C_t) \longmapsto \text{cost of } C_t$

Figure: Real-time segmentation (left) without and (right) with p2p-cleaning

Applications (5)

back

Modifying MA

- add new poses \longmapsto real-time segmentation

example

Applications (5)

back

Modifying MA

- add new poses \mapsto real-time segmentation
- edit arbitrary pose p_t : [replace] $C_t \mapsto C'_t$ & [recompute OS]

example

example

Applications (5)

Modifying MA

- add new poses \mapsto real-time segmentation
- edit arbitrary pose p_t : [replace] $C_t \mapsto C'_t$ & [recompute OS]

example

example

Segmentation transfer

- MA 1
 - MA 2
 - ...
 - MA k
- } [merge MAs] \mapsto [merge segmentations]

example

Experimental Study (1)

Performance Analysis (1)

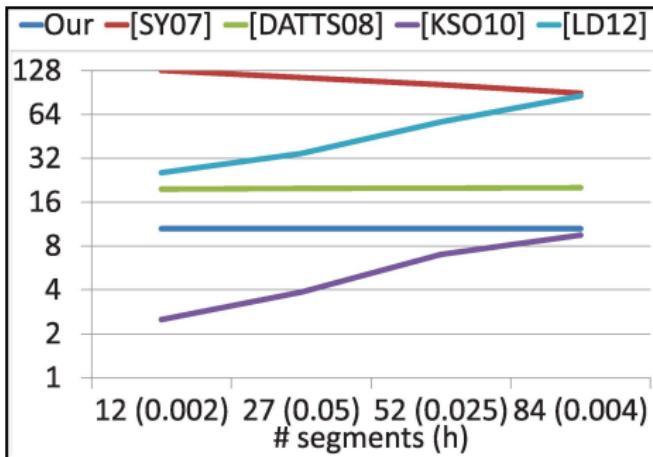
- linear behavior on the geometry size
- linear scale when increasing per-pose partitionings
- pose decompositions consumes most of the time
- interactive performance \longleftarrow [GPU-accelerated clustering]

Mesh Animation	Per pose								Mesh Animation				
	Vertices	Faces	Poses	Clusters	Feature	Clustering		Variable Segmentation		Over-segmentation	p2p-Cleaning	Total	
						Compute	Propagate Colors	Compute	Propagate Colors			Compute (segments)	Compute (segments)
Mesh Animation	Vertices	Faces	Poses	Clusters	Feature	Compute	Propagate Colors	Compute	Propagate Colors	Over-segmentation	p2p-Cleaning	real-time	off-line
Hand	7929	15855	22	12	0.116	1.603	0.0014	0.0046	0.00084	0.215 (379)	0.03 (24)	1.72	37.95
Elephant Gallop	42321	84638	24	5	0.483	3.327	0.0053	0.012	0.0041	0.735 (743)	0.07 (18)	3.844	92.25
				10	0.483	6.463	0.0053	0.012	0.0041	1.475 (1729)	0.09 (18)	7.281	174.7
Flowing Cloth	25921	51200	19	[2,5]	0.295	1.873	0.0028	0.0074	0.0025	0.698 (400)	0.02 (24)	2.205	41.91
Samba	9971	19938	24	5	0.158	0.734	0.0013	0.024	0.0009	0.78 (1016)	0.07 (12)	0.927	22.26

Experimental Study (1)

Performance Analysis (2)

- tiny cost when moving from high to low resolutions
- only multi-source region growing method¹ is faster



¹ [Kavan et al., CGF'10]

Experimental Study (2)

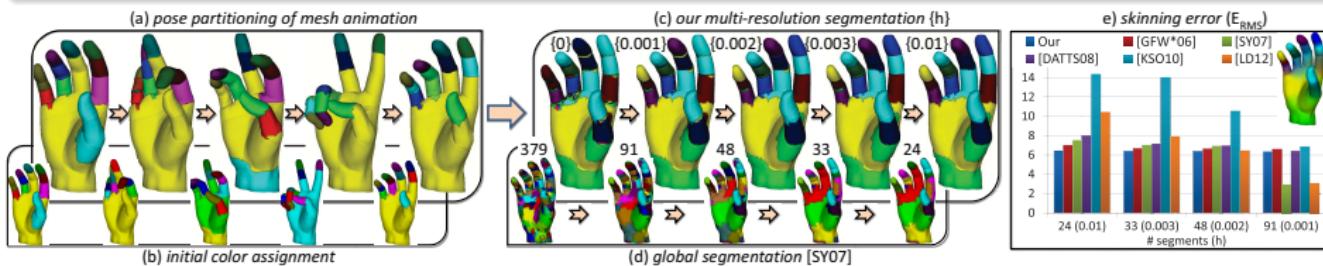
Quality Analysis (1)

- E_{RMS} : skinning approximation error
 - [rigid] \longmapsto [volumetric parts]
 - [highly-deformable] \longmapsto [surface parts]
- } [hybrid]

Experimental Study (2)

Quality Analysis (1)

- E_{RMS} : skinning approximation error
 - [rigid] \longmapsto [volumetric parts]
 - [highly-deformable] \longmapsto [surface parts]
- } [hybrid]

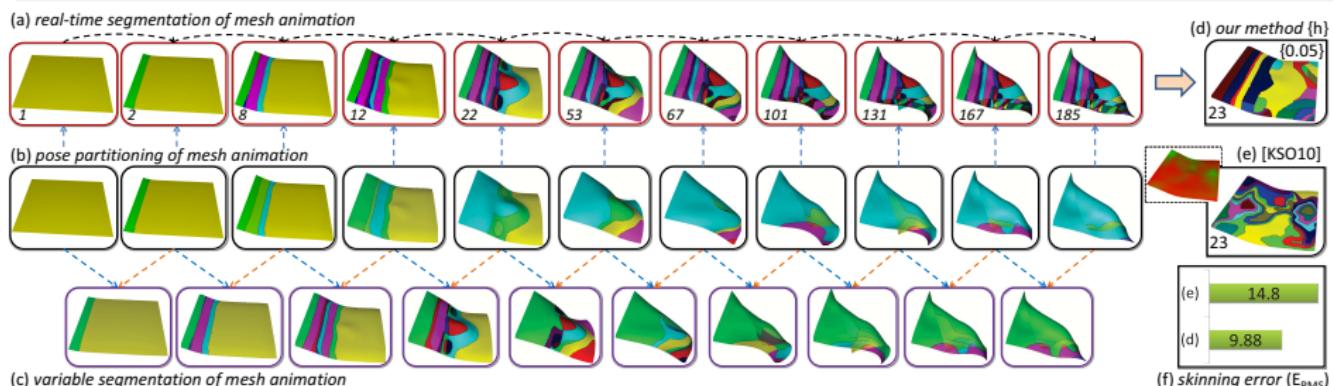


¹[Kavan et al., CGF'10], ²[Schaeffer et al., GP'07]

Experimental Study (2)

Quality Analysis (1)

- E_{RMS} : skinning approximation error
 - [rigid] \longmapsto [volumetric parts]
 - [highly-deformable] \longmapsto [surface parts]
- } [hybrid]



¹[Kavan et al., CGF'10], ²[Schaeffer et al., GP'07]

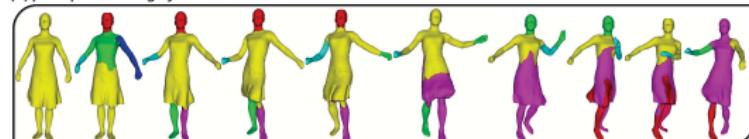
Experimental Study (2)

Quality Analysis (1)

- E_{RMS} : skinning approximation error
- [rigid] \longmapsto [volumetric parts]
- [highly-deformable] \longmapsto [surface parts]

} [hybrid]

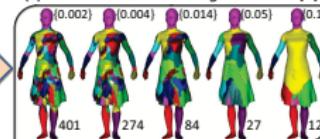
(a) pose partitioning of mesh animation



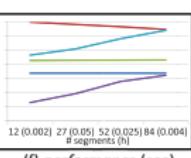
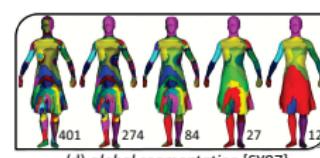
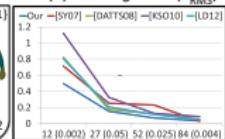
(b) variable segmentation of mesh animation

401	274	84	27	12
-----	-----	----	----	----

(c) our multi-resolution segmentation {h}



(e) skinning error (E_{RMS})



(d) global segmentation [SY07]

(f) performance (sec)

¹[Kavan et al., CGF'10], ²[Schaeffer et al., GP'07]

Discussion

Quality

- [our method] (merge & simplify) VS [global segmentation]

Discussion

Quality

- [our method] (merge & simplify) VS [global segmentation]
 - depends on quality/number of individual pose partitionings:
 - ① higher-quality clustering
 - ② boundary smoothing
 - ③ confidence measures
- } parameter tuning & extra computation

Discussion

Quality

- [our method] (merge & simplify) VS [global segmentation]
 - depends on quality/number of individual pose partitionings:
 - ① higher-quality clustering
 - ② boundary smoothing
 - ③ confidence measures
- } parameter tuning & extra computation

Performance

- applications run at real-time speed

Discussion

Quality

- [our method] (merge & simplify) VS [global segmentation]
 - depends on quality/number of individual pose partitionings:
 - ① higher-quality clustering
 - ② boundary smoothing
 - ③ confidence measures
- } parameter tuning & extra computation

Performance

- applications run at real-time speed
- reduce computation cost of pose partitioning:
 - ① cluster poses with high partitioning similarity

Discussion

Quality

- [our method] (merge & simplify) VS [global segmentation]
 - depends on quality/number of individual pose partitionings:
 - ① higher-quality clustering
 - ② boundary smoothing
 - ③ confidence measures
- } parameter tuning & extra computation

Performance

- applications run at real-time speed
- reduce computation cost of pose partitioning:
 - ① cluster poses with high partitioning similarity
 - ② GPU-accelerated clustering

Outline

- 1 Background and Preliminaries
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 S-buffer: Sparsity-aware multi-fragment rendering
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 k^+ -buffer: Fragment synchronized k -buffer
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Pose-to-pose skinning of MAs

Skinning MAs

[go back](#)

- [key-frame compression]: $3n \cdot k \longmapsto (3n + 4n) + y \cdot k, y = \{8, 12\}$

Pose-to-pose skinning of MAs

Skinning MAs

[go back](#)

- [key-frame compression]: $3n \cdot k \longmapsto (3n + 4n) + y \cdot k, y = \{8, 12\}$
- [additional applications]: GPU reproduction, lod selection, etc

Pose-to-pose skinning of MAs

go back

Skinning MAs

- [key-frame compression]: $3n \cdot k \longmapsto (3n + 4n) + y \cdot k, y = \{8, 12\}$
- [additional applications]: GPU reproduction, lod selection, etc

Our contribution - *p2p-skinning*[†]

- ① [Idea:] $p_{t-1}^s \longmapsto p_t^a$ instead of $p_0 \longmapsto p_t^a$ - (reproduction scheme)

[†]A. Vasilakis & G. Antonopoulos & I. Fudos, *Pose-to-Pose Skinning of Animated Meshes*, Symposium on Computer Animation (Posters), Vancouver, Canada, August 5-7, 2011.

Pose-to-pose skinning of MAs

go back

Skinning MAs

- [key-frame compression]: $3n \cdot k \longmapsto (3n + 4n) + y \cdot k, y = \{8, 12\}$
- [additional applications]: GPU reproduction, lod selection, etc

Our contribution - *p2p-skinning*[†]

- ① [Idea:] $p_{t-1}^s \longmapsto p_t^a$ instead of $p_0 \longmapsto p_t^a$ - (reproduction scheme)
- ② supports both [linear] (LBS) and [non-linear] (DQS) skinning

[†]A. Vasilakis & G. Antonopoulos & I. Fudos, *Pose-to-Pose Skinning of Animated Meshes*, Symposium on Computer Animation (Posters), Vancouver, Canada, August 5-7, 2011.

Pose-to-pose skinning of MAs

go back

Skinning MAs

- [key-frame compression]: $3n \cdot k \longmapsto (3n + 4n) + y \cdot k, y = \{8, 12\}$
- [additional applications]: GPU reproduction, lod selection, etc

Our contribution - *p2p-skinning*[†]

- ① [Idea:] $p_{t-1}^s \longmapsto p_t^a$ instead of $p_0 \longmapsto p_t^a$ - (reproduction scheme)
- ② supports both [linear] (LBS) and [non-linear] (DQS) skinning
- ③ enables prior applications + [arbitrary pose editing]

[†]A. Vasilakis & G. Antonopoulos & I. Fudos, *Pose-to-Pose Skinning of Animated Meshes*, Symposium on Computer Animation (Posters), Vancouver, Canada, August 5-7, 2011.

Pose-to-pose skinning of MAs

go back

Skinning MAs

- [key-frame compression]: $3n \cdot k \longmapsto (3n + 4n) + y \cdot k, y = \{8, 12\}$
- [additional applications]: GPU reproduction, lod selection, etc

Our contribution - *p2p-skinning*[†]

- ① [Idea:] $p_{t-1}^s \longmapsto p_t^a$ instead of $p_0 \longmapsto p_t^a$ - (reproduction scheme)
- ② supports both [linear] (LBS) and [non-linear] (DQS) skinning
- ③ enables prior applications + [arbitrary pose editing]
- ④ offers [refinements] to improve skinning quality

[†]A. Vasilakis & G. Antonopoulos & I. Fudos, *Pose-to-Pose Skinning of Animated Meshes*, Symposium on Computer Animation (Posters), Vancouver, Canada, August 5-7, 2011.

Approximating MAs with Skinning

Notation & Equations

- [rest-pose]: $\{v_j^0\}$
 - [weights]: $[w_{1,j}, \dots, w_{B,j}]$, $w_{i,j} > 0$, $\sum_{i=1}^B w_{i,j} = 1$
 - [transformation matrices]: $[M_1^t, \dots, M_B^t]$ - **unknown**
- $\left. \begin{array}{l} \\ \\ \end{array} \right\}$ known

$$v_j^t = T_j^t \cdot v_j^0, \forall t \in [0, k], \text{ where } T_j^t = \sum_{b=1}^B w_{b,j} M_b^t \quad (1)$$

Approximating MAs with Skinning

Notation & Equations

- [rest-pose]: $\{v_j^0\}$
- [weights]: $[w_{1,j}, \dots, w_{B,j}]$, $w_{i,j} > 0$, $\sum_{i=1}^B w_{i,j} = 1$
- [transformation matrices]: $[M_1^t, \dots, M_B^t]$ - **unknown**

$$v_j^t = T_j^t \cdot v_j^0, \forall t \in [0, k], \text{ where } T_j^t = \sum_{b=1}^B w_{b,j} M_b^t \quad (1)$$

- [dual quaternions]: less storage, only rigid motion.

Approximating MAs with Skinning

Notation & Equations

- [rest-pose]: $\{v_j^0\}$
- [weights]: $[w_{1,j}, \dots, w_{B,j}]$, $w_{i,j} > 0$, $\sum_{i=1}^B w_{i,j} = 1$
- [transformation matrices]: $[M_1^t, \dots, M_B^t]$ - **unknown**

$$v_j^t = T_j^t \cdot v_j^0, \forall t \in [0, k], \text{ where } T_j^t = \sum_{b=1}^B w_{b,j} M_b^t \quad (1)$$

- [dual quaternions]: less storage, only rigid motion.

Fitting Formula

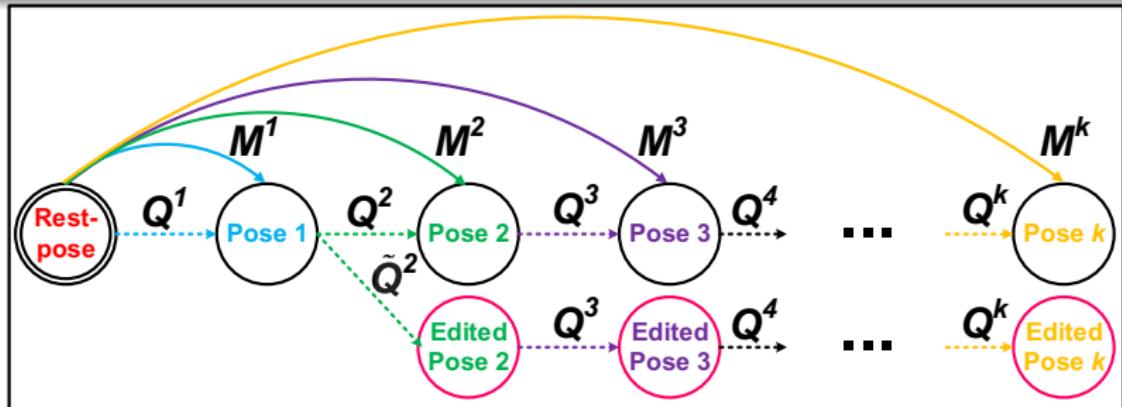
$$\min \left\{ \sum_{j=0}^{n-1} \| \mathbf{v}_j^t - v_j^t \|^2 \right\}, \text{ where } \mathbf{v}_j^t \in p_t$$

Pose-to-pose Fitting

Our solution

- small deformation variations will occur between sequential poses
- reformulate equation (1):

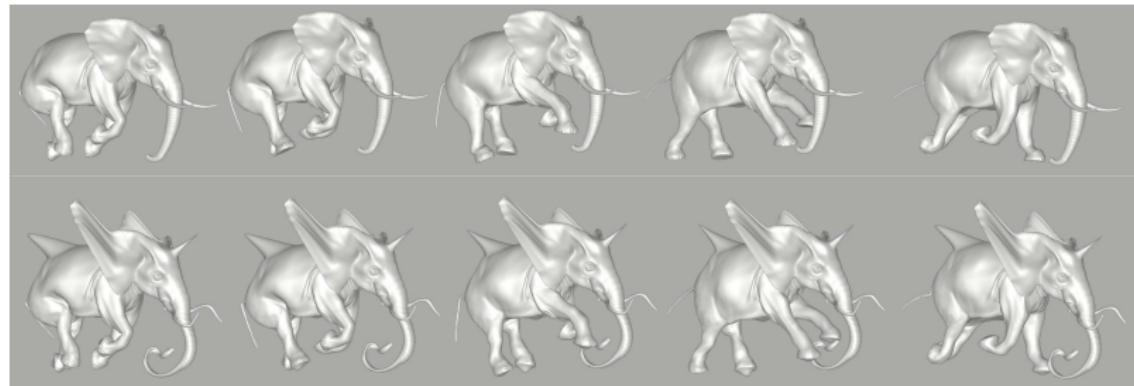
$$v_j^t = \left(\sum_{b=1}^B w_{b,j} Q_b^t \right) v_j^{t-1}, \quad T_j^t = \left(\sum_{b=1}^B w_{b,j} Q_b^t \right) T_j^{t-1} \quad (2)$$



Animation Editing

Rest-pose editing

- simply modify $\{v_j^0\}$

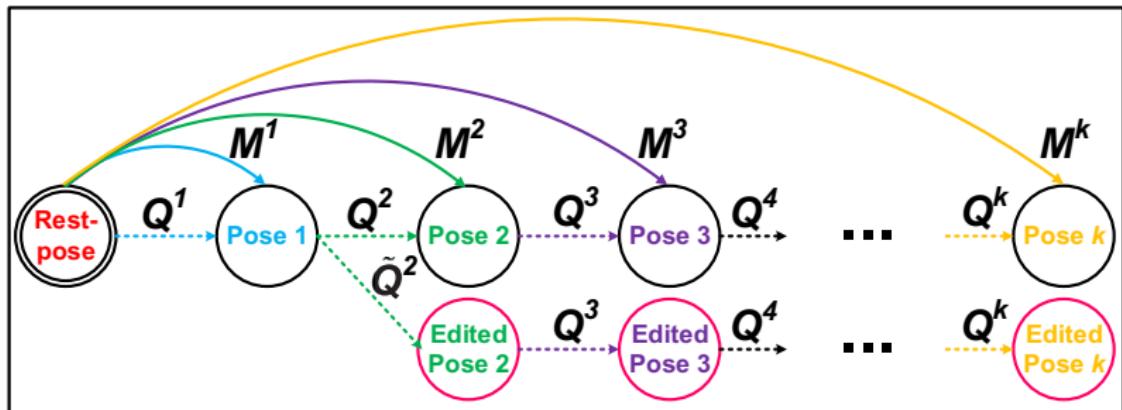


Animation Editing

Arbitrary pose editing

- $\{\tilde{Q}_b^e\}$ matrices: $p_{e-1} \longmapsto p_e \Rightarrow$

$$T_j^t = T_j^{t,e-1} \left(\sum_{b=1}^B w_{b,j} \tilde{Q}_b^e \right) T_i^{e-1,1}, \quad T_j^{t_0,t_1} = \prod_{t=t_0}^{t_1} \left(\sum_{b=1}^B w_{b,j} Q_b^t \right)$$

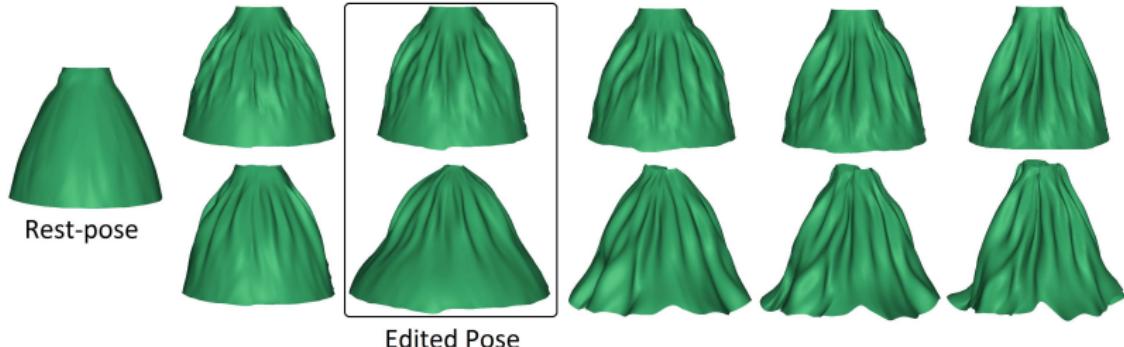


Animation Editing

Arbitrary pose editing

- $\{\tilde{Q}_b^e\}$ matrices: $p_{e-1} \longmapsto p_e \Rightarrow$

$$T_j^t = T_j^{t,e-1} \left(\sum_{b=1}^B w_{b,j} \tilde{Q}_b^e \right) T_i^{e-1,1}, \quad T_j^{t_0,t_1} = \prod_{t=t_0}^{t_1} \left(\sum_{b=1}^B w_{b,j} Q_b^t \right)$$

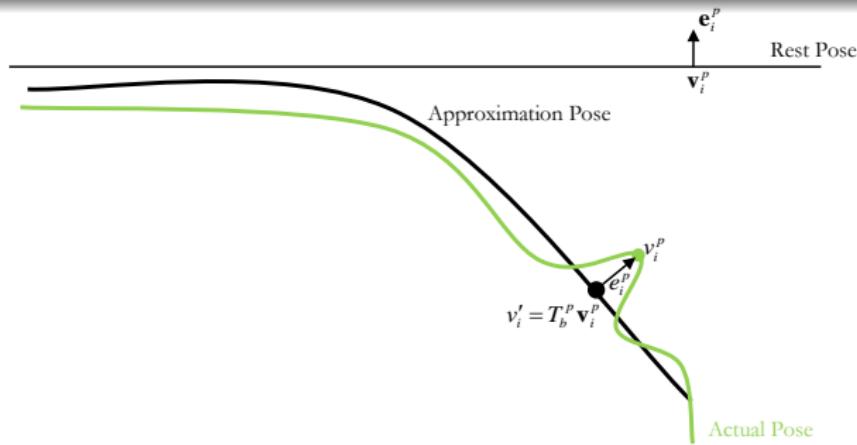


Rest-pose Vertex Corrections

Problem formulation

- [displacement field]: $e^V = [e_0^V, \dots, e_{n-1}^V] \in \mathbb{R}^{4 \times n}$:

$$\min \left\{ \sum_{t=0}^{k-1} \left\| T^t(v_j^0 + e_j^V) - \mathbf{v}_j^t \right\|^2 \right\}, \quad j \in [0, \dots, n] \quad (3)$$



Weight Corrections

Problem formulation

- [displacement field]: $e^W = [e_1^W, \dots, e_B^W] \in \mathbb{R}^{1 \times B}$

$$\min \left\{ \sum_{t=0}^{k-1} \left\| \sum_{b=1}^B (w_{b,j} + e_{b,j}^W) Q_b^t \mathbf{v}_j^{t-1} - \mathbf{v}_j^t \right\|^2 \right\}, \quad j \in [0, \dots, n] \quad (4)$$

Weight Corrections

Problem formulation

- [displacement field]: $e^W = [e_1^W, \dots, e_B^W] \in \mathbb{R}^{1 \times B}$

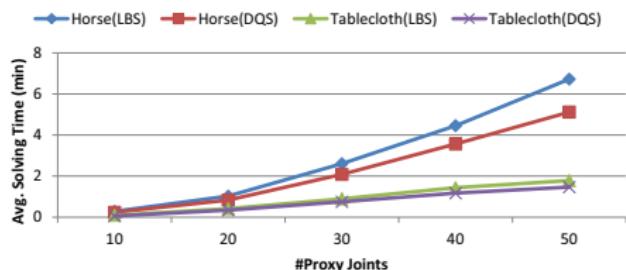
$$\min \left\{ \sum_{t=0}^{k-1} \left\| \sum_{b=1}^B (w_{b,j} + e_{b,j}^W) Q_b^t \mathbf{v}_j^{t-1} - \mathbf{v}_j^t \right\|^2 \right\}, \quad j \in [0, \dots, n] \quad (4)$$

Avoid Over-fitting

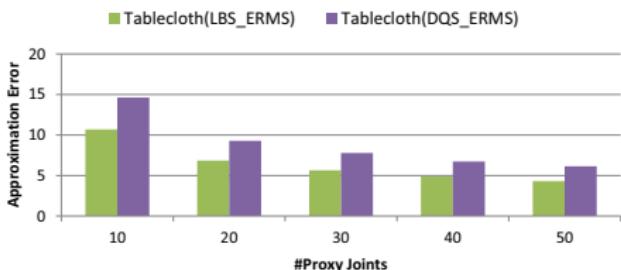
- [convexity constraint]:
 $\sum_{b=1}^B (w_{b,j} + e_{b,j}^W) = 1 \Leftrightarrow \sum_{b=1}^{B-1} e_{b,j}^W = -e_{B,j}^W, \quad j \in [0, \dots, n]$
- [non-negativity constraint]:
 - $w_{b,j} + e_{b,j}^W \geq 0 \Leftrightarrow e_{b,j}^W \geq -w_{b,j}, \quad b \in [1, \dots, B-1]$
 - $w_{B,j} + e_{B,j}^W \geq 0 \Leftrightarrow \sum_{b=1}^{B-1} e_{b,j}^W \leq w_{B,j}$

Performance/Quality Evaluation (LBS vs DQS)

LBS vs DQS Avg. Solving Time



LBS vs DQS Fitting Error



Performance/Quality Evaluation (P2P-Skin vs Prior Art)

Source Data				Skinning Methods					
				SAD		FESAM		p2p-skinned	
Name	Vertices	Poses	Proxy-Joints	Error	Time	Error	Time	Error	Time
Elephant	42321	48	25	13.3	368.1	1.39	778.9	2.2	420.1
Samba	9971	175	30	12.6	380	1.17	344.2	1.98	443.5
Skirt	5095	60	75	4.8	422.6	1.72	264.9	2.81	550.9
Facial Expression	23725	23	50	38.7	283.8	5.2	551.1	7.5	332.1

¹ SAD: [Kavan et al. I3D'07], ² FESAM:[Kavan et al. CGF'10]

Performance/Quality Evaluation (P2P-Skin vs Prior Art)



Figure: (From left to right) An approximated pose using SAD¹, using FESAM², using p2p-skinning, using p2p-skinning with corrections and finally the original pose

¹ SAD: [Kavan et al. I3D'07], ² FESAM:[Kavan et al. CGF'10]

Conclusions

Approximating MAs with pose-to-pose skinning

- handles both linear and non-linear fitting
- supports rigid, highly-deformable and hybrid MAs
- enables pose editing of arbitrary animation frames
- increases skinning quality via novel vertex and weight corrections

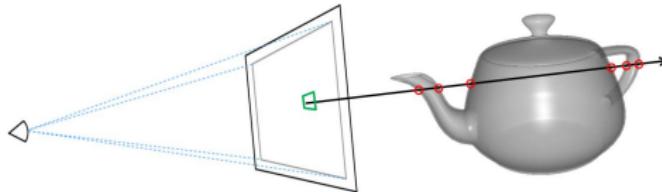
Outline

- 1 Background and Preliminaries
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 **S-buffer: Sparsity-aware multi-fragment rendering**
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 k^+ -buffer: Fragment synchronized k -buffer
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Revisiting A-buffer architecture

Memory-friendly Multi-fragment Rendering

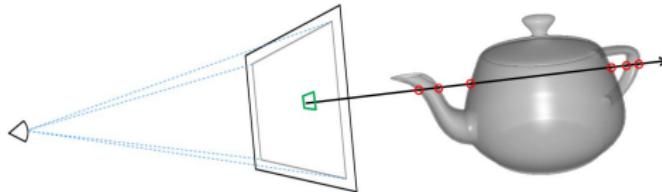
- capture all generated fragments
- memory overflow & performance bottlenecks



Revisiting A-buffer architecture

Memory-friendly Multi-fragment Rendering

- capture all generated fragments
- memory overflow & performance bottlenecks



Our contribution - *S-buffer*[†]

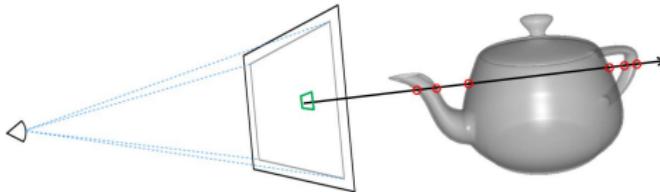
- ① precise memory allocation

[†]A. Vasilakis & I. Fudos, *S-buffer: Sparsity-aware Multi-fragment Rendering*, Eurographics 2012 (Short Papers), Cagliari, Italy, May 13-18, 2012.

Revisiting A-buffer architecture

Memory-friendly Multi-fragment Rendering

- capture all generated fragments
- memory overflow & performance bottlenecks



Our contribution - *S-buffer*[†]

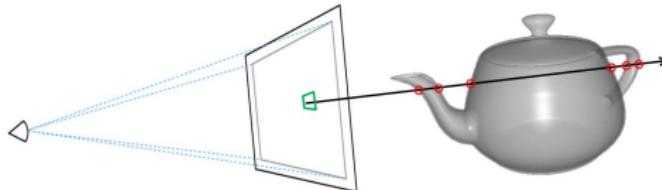
- ① precise memory allocation
- ② reduces fragment & pixel contention

[†]A. Vasilakis & I. Fudos, *S-buffer: Sparsity-aware Multi-fragment Rendering*, Eurographics 2012 (Short Papers), Cagliari, Italy, May 13-18, 2012.

Revisiting A-buffer architecture

Memory-friendly Multi-fragment Rendering

- capture all generated fragments
- memory overflow & performance bottlenecks



Our contribution - *S-buffer*[†]

- ① precise memory allocation
- ② reduces fragment & pixel contention
- ③ two-rendering-pass method

[†]A. Vasilakis & I. Fudos, *S-buffer: Sparsity-aware Multi-fragment Rendering*, Eurographics 2012 (Short Papers), Cagliari, Italy, May 13-18, 2012.

Related Work

Geometry sorting methods

- object sorting
- primitive sorting

Related Work

Geometry sorting methods

- object sorting
- primitive sorting

Fragment sorting methods

- depth-peeling

Related Work

Geometry sorting methods

- object sorting
- primitive sorting

Fragment sorting methods

- depth-peeling
- hardware-implemented buffers

Related Work

Geometry sorting methods

- object sorting
- primitive sorting

Fragment sorting methods

- depth-peeling
- hardware-implemented buffers

Framework design goals

Related Work

Geometry sorting methods

- object sorting
- primitive sorting

Fragment sorting methods

- depth-peeling
- hardware-implemented buffers

Framework design goals

- Quality: Fragment extraction accuracy (**A**)

Related Work

Geometry sorting methods

- object sorting
- primitive sorting

Fragment sorting methods

- depth-peeling
- hardware-implemented buffers

Framework design goals

- Quality: Fragment extraction accuracy (**A**)
- Time performance (**P**)

Related Work

Geometry sorting methods

- object sorting
- primitive sorting

Fragment sorting methods

- depth-peeling
- hardware-implemented buffers

Framework design goals

- Quality: Fragment extraction accuracy (**A**)
- Time performance (**P**)
- Memory allocation (**Ma**) and caching (**Mc**)

Related Work

Geometry sorting methods

- object sorting
- primitive sorting

Fragment sorting methods

- depth-peeling
- hardware-implemented buffers

Framework design goals

- Quality: Fragment extraction accuracy (**A**)
- Time performance (**P**)
- Memory allocation (**Ma**) and caching (**Mc**)
- GPU capabilities (**G**)

Related Work - Depth Peeling

Methods

- Front-to-back (F2B)¹

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, HPG'09]

Related Work - Depth Peeling

Methods

- Front-to-back (F2B)¹
- Dual direction (DUAL)²

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, HPG'09]

Related Work - Depth Peeling

Methods

- Front-to-back (F2B)¹
- Dual direction (DUAL)²
- Uniform bucket (BUN)³

Characteristics

- A: z-fighting [↓] - P: multi-pass [↓]

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, HPG'09]

Related Work - Depth Peeling

Methods

- Front-to-back (F2B)¹
- Dual direction (DUAL)²
- Uniform bucket (BUN)³

Characteristics

- A: z-fighting [↓] - P: multi-pass [↓]
- Ma: low [↑] - Mc: fast [↑]

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, HPG'09]

Related Work - Depth Peeling

Methods

- Front-to-back (F2B)¹
- Dual direction (DUAL)²
- Uniform bucket (BUN)³

Characteristics

- A: z-fighting [↓] - P: multi-pass [↓]
- Ma: low [↑] - Mc: fast [↑]
- G: old & modern [↑]

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, HPG'09]

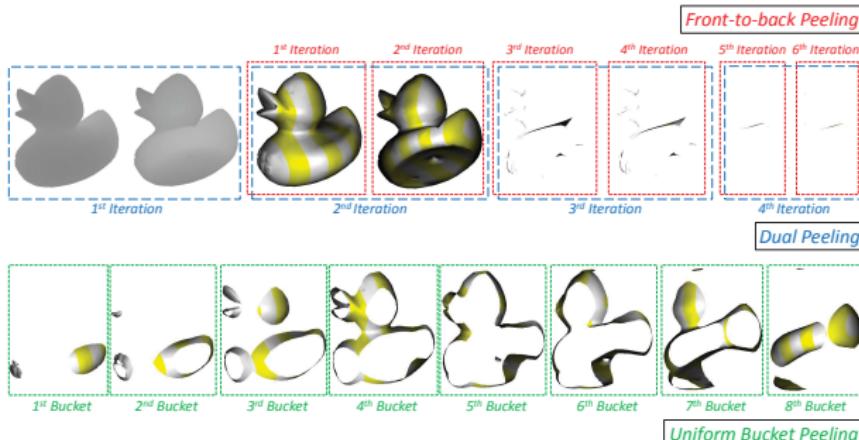
Related Work - Depth Peeling

Methods

- Front-to-back (F2B)¹
- Dual direction (DUAL)²
- Uniform bucket (BUN)³

Characteristics

- A: z-fighting [↓] - P: multi-pass [↓]
- Ma: low [↑] - Mc: fast [↑]
- G: old & modern [↑]

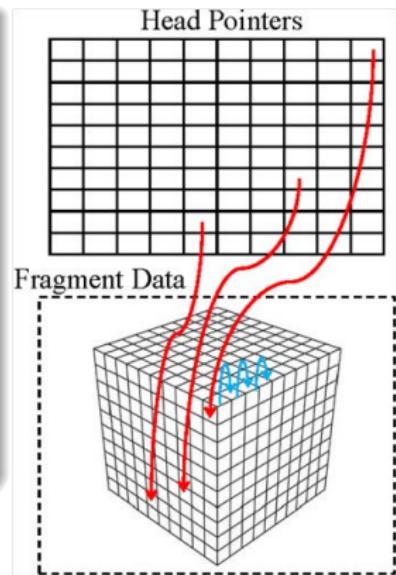


¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, HPG'09]

Related Work - Hardware Buffers (1)

Per-pixel Fixed-size Arrays

- **M**: huge (unused) [↓]
- **M**: very fast [↑]
- **G**: [commodity]: KB¹, KB-SR²
 - **A**: RMW hazards - 8 fragments [↓]
 - **P**: fast [↑]
- **G**: [modern]: AB_{FP}^{3,4}
 - **A**: 100% if enough memory [↑]
 - **P**: fastest (single pass) [↑]

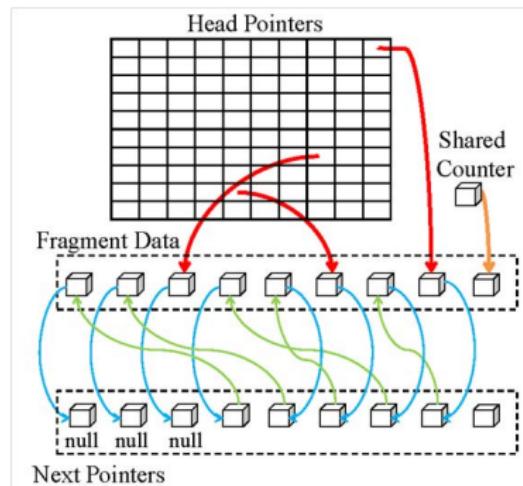


¹ [Bavoil et al, I3D'07], ² [Myers et al, SIGGRAPH'07], ³ [Liu et al, I3D'10], ⁴ [Crassin, blog'11]

Related Work - Hardware Buffers (2)

Per-pixel Linked Lists (AB_{LL})¹

- A: 100% if enough memory [\uparrow]
- P: fast (fragment contention) [\downarrow]
- Mc: low cache hit ratio [\downarrow]
- Ma: high
 - if [overflow]: accurate reallocation (extra pass needed) [\downarrow]
 - wasted memory [\downarrow]
- G: modern cards [\downarrow]

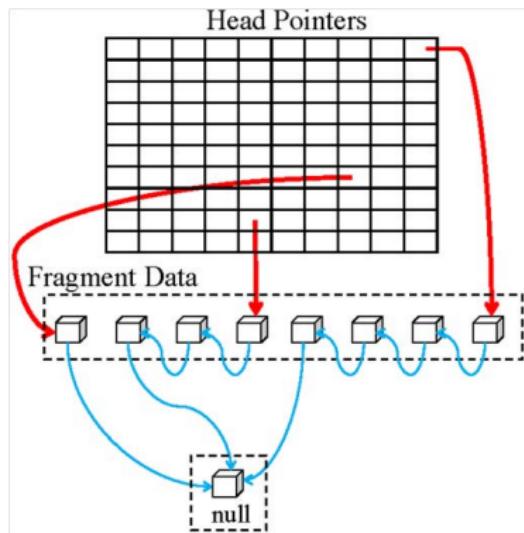


¹ [Yang et al, CGF'10]

Related Work - Hardware Buffers (3)

Per-pixel Variable-length Arrays

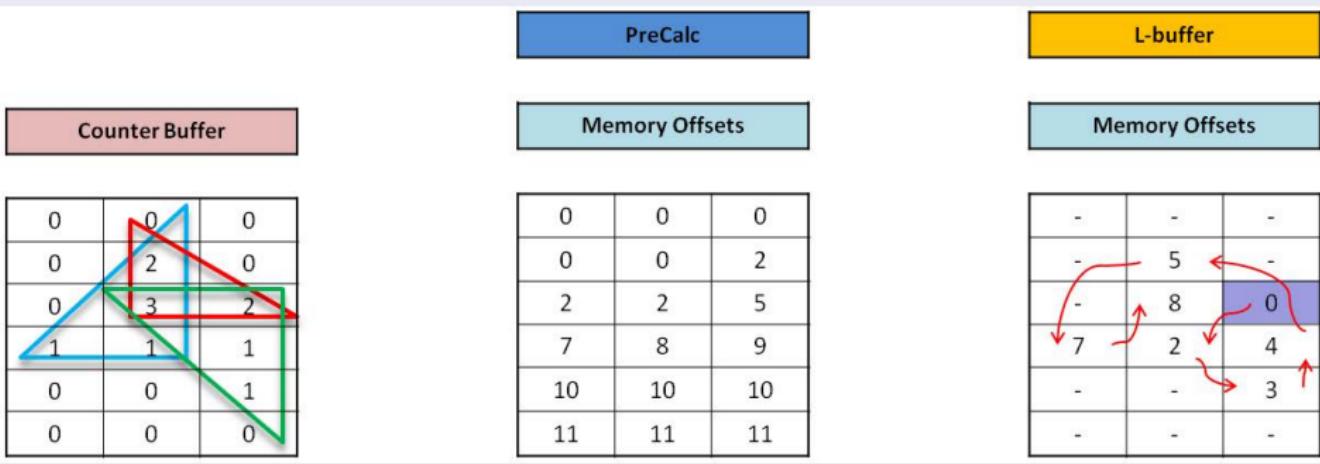
- A: 100% if enough memory [↑]
- P: fast (2 passes needed) [↑]
- Ma: precise allocation [↑]
- Mc: fast [↑]
- G: [commodity]
 - PreCalc¹ (prefix sum)
 - I-buffer² (randomized prefix sum)
- G: [modern]
 - DFB³ (parallel prefix sum)



¹ [Peeper, patent'08], ² [Lipowski, ICCVG'10], ³ [Maule et al, CG'14]

Related Work - Hardware Buffers (3)

Example: PreCalc¹ & I-buffer²



¹ [Peeper, patent'08], ² [Lipowski, ICCVG'10]

S-buffer Pipeline

1. Fragment Count Rendering Pass

- # fragments per pixel
- total generated fragments

S-buffer Pipeline

2. Memory Referencing Full-screen Pass

[parallelized randomized prefix sum]:

- S multiple [shared counters]: $C = \{C(0), \dots, C(S - 1)\}$

S-buffer Pipeline

2. Memory Referencing Full-screen Pass

[parallelized randomized prefix sum]:

- S multiple [shared counters]: $C = \{C(0), \dots, C(S - 1)\}$
- simple [hash function]: $H(p) = (p.x + sc.width * p.y) \% S$

S-buffer Pipeline

2. Memory Referencing Full-screen Pass

[parallelized randomized prefix sum]:

- S multiple [shared counters]: $C = \{C(0), \dots, C(S - 1)\}$
- simple [hash function]: $H(p) = (p.x + sc.width * p.y) \% S$
- [sequential prefix sum] on shared counters: $C_{pr}(i) = \sum_0^{i-1} C(i)$

S-buffer Pipeline

2. Memory Referencing Full-screen Pass

[parallelized randomized prefix sum]:

- S multiple [shared counters]: $C = \{C(0), \dots, C(S-1)\}$
- simple [hash function]: $H(p) = (p.x + sc.width * p.y) \% S$
- [sequential prefix sum] on shared counters: $C_{pr}(i) = \sum_0^{i-1} C(i)$
- [inverse mapping extension]:
 - [split] C : $G_1 = \{C(0), \dots, C(\lfloor \frac{S}{2} \rfloor)\}, G_2 = \{C(\lfloor \frac{S}{2} \rfloor + 1), \dots, C(S-1)\}$

S-buffer Pipeline

2. Memory Referencing Full-screen Pass

[parallelized randomized prefix sum]:

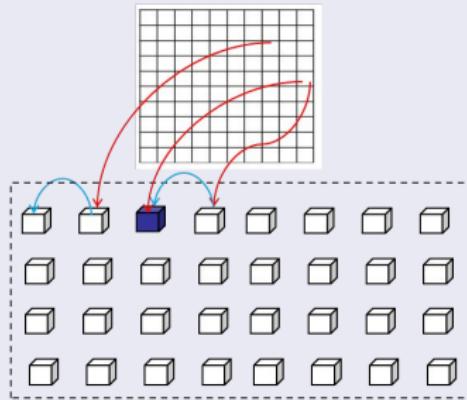
- S multiple [shared counters]: $C = \{C(0), \dots, C(S-1)\}$
- simple [hash function]: $H(p) = (p.x + sc.width * p.y) \% S$
- [sequential prefix sum] on shared counters: $C_{pr}(i) = \sum_0^{i-1} C(i)$
- [inverse mapping extension]:
 - [split] C : $G_1 = \{C(0), \dots, C(\lfloor \frac{S}{2} \rfloor)\}, G_2 = \{C(\lfloor \frac{S}{2} \rfloor + 1), \dots, C(S-1)\}$
 - [final memory offset]:

$$p.\text{offset} = \begin{cases} A(p), & \text{if } p \in G_1 \\ \text{total_fragments-1-}A(p), & \text{otherwise} \end{cases}$$

where $A(p) = p.\text{address} + C_{pr}(H(p))$

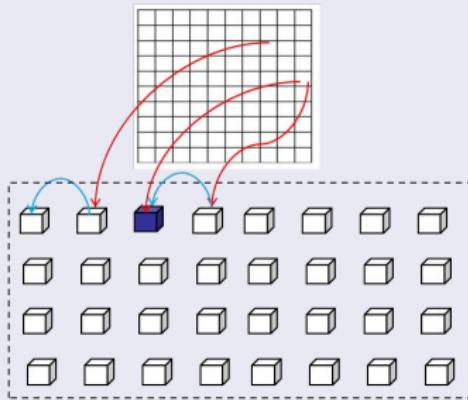
S-buffer Pipeline

3. Fragment Storing Rendering Pass



S-buffer Pipeline

3. Fragment Storing Rendering Pass

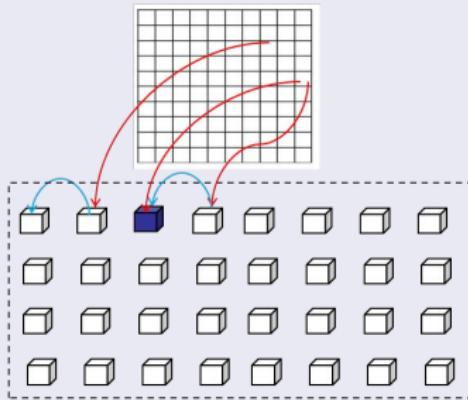


4. Resolve Full-screen Pass

- ① [Fragment sorting]
- ② [Effect]

S-buffer Pipeline

3. Fragment Storing Rendering Pass



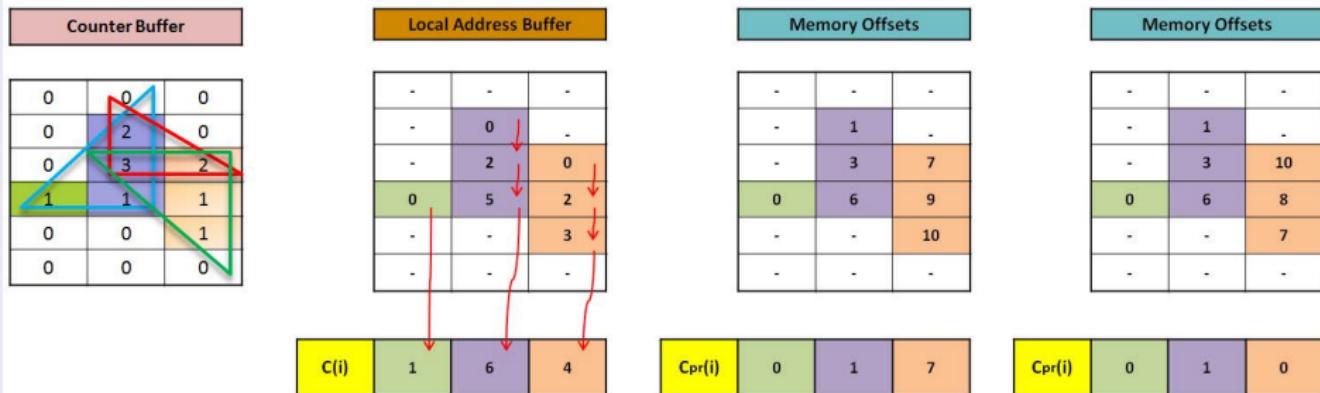
4. Resolve Full-screen Pass

- ① [Fragment sorting]
- ② [Effect]

S-buffer Pipeline

Example: using 3 shared counters

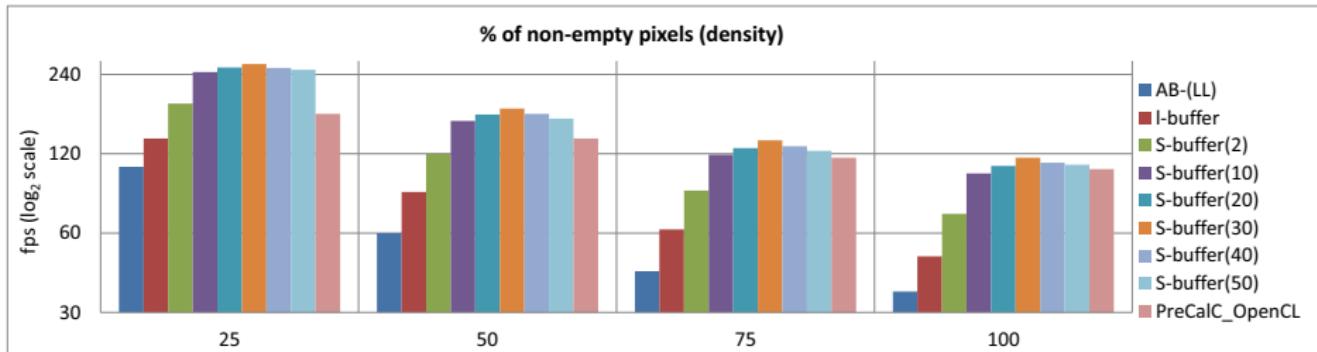
Inverse mapping



Performance Analysis (1)

Impact of Sparsity

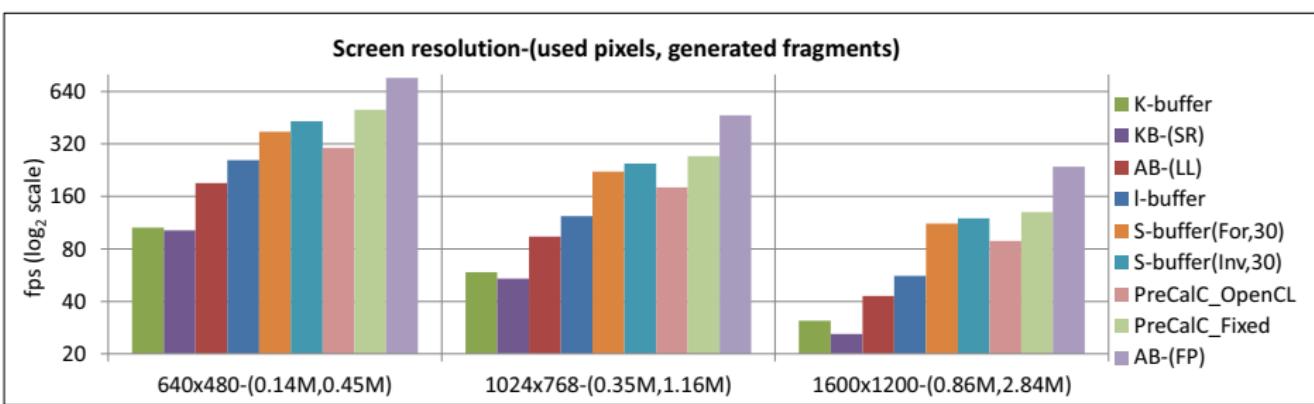
- [Scene]: 70,000 triangles, 12 layers, 1024^2 viewport
- $[AB_{LL}]$: $O(m)$, $m(>n) = \text{total fragments}$
- [I-buffer]: $O(n)$, $n = \text{non-empty pixels}$
- [S-buffer] speeds up: n/S , S shared counters
- [PreCalc_OpenCL]: OpenGL/OpenCL syncing time



Performance Analysis (2)

Impact of Resolution

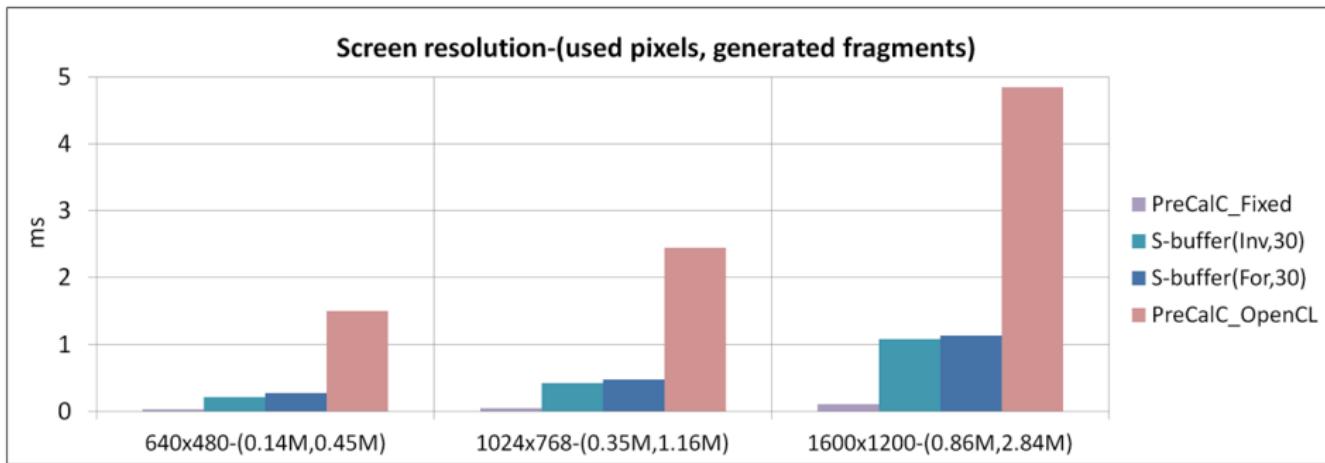
- [Scene]: 110,000 triangles, 25 layers, $p_d = 45\%$
- [S-buffer] = 85% of PreCalc_Fixed speed
- [forward] vs [inverse mapping]



Performance Analysis (2)

Impact of Resolution

- [Scene]: 110,000 triangles, 25 layers, $p_d = 45\%$
- [S-buffer] = 85% of PreCalc_Fixed speed
- [forward] vs [inverse mapping]

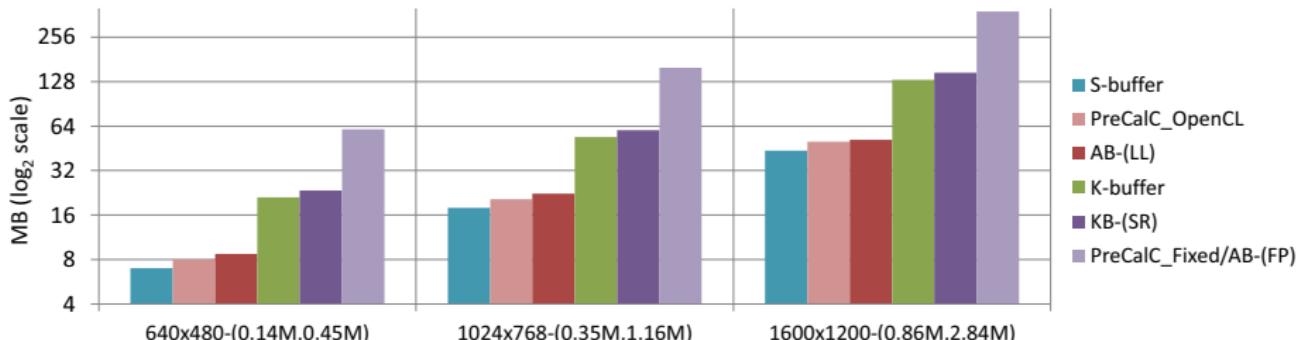


Memory Allocation Analysis

Impact of Resolution

- [Scene]: 110,000 triangles, 25 layers, $p_d = 45\%$
- [fixed-sized arrays]:
 - AB_{FP}: Wasted resources (88%)
 - KB,KB-SR: 30% less memory due to 8 fragments/pixel
- [AB_{LL}]: extra memory (pointer list)

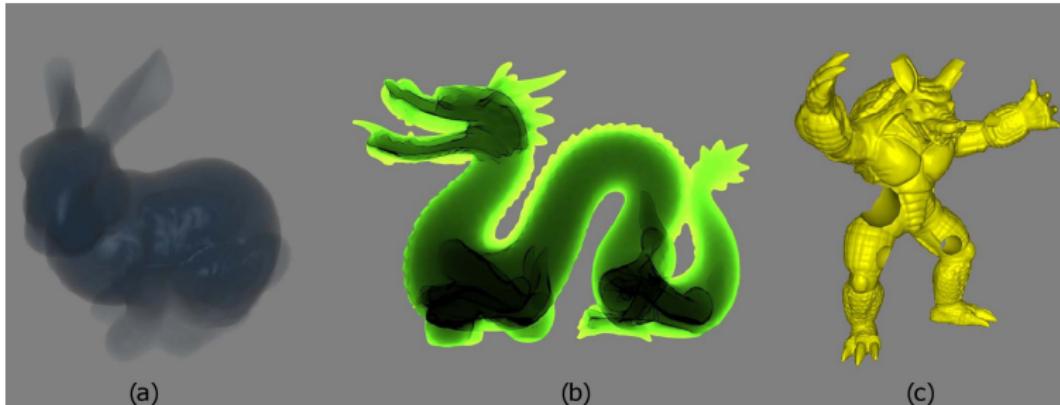
Screen resolution-(used pixels, generated fragments)



Conclusions

Multi-fragment rendering via S-buffer

- two-geometry-pass GPU-accelerated A-buffer that take advantage of
 - the fragment distribution (exact memory allocation)
 - the sparsity of the pixel-space (improved performance)
- organizes storage into variable contiguous regions for each pixel
- integrates into the standard graphics pipeline



Outline

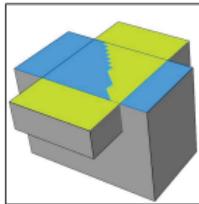
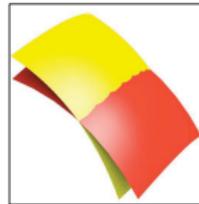
- 1 Background and Preliminaries
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 S-buffer: Sparsity-aware multi-fragment rendering
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 k^+ -buffer: Fragment synchronized k -buffer
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Extracting coplanar fragments

Z-fighting phenomenon

more

- two or more primitives have same depth
- intersecting or overlapping surfaces
- multifragment rasterization is even more susceptible

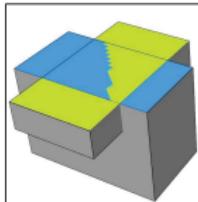
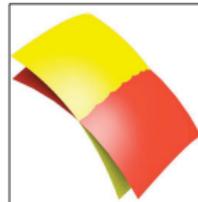


Extracting coplanar fragments

Z-fighting phenomenon

more

- two or more primitives have same depth
- intersecting or overlapping surfaces
- multifragment rasterization is even more susceptible



Our contribution[†]

- ① [robust] & [approximate] methods

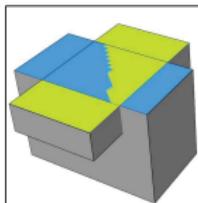
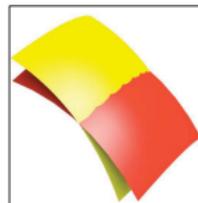
[†]A. Vasilakis & I. Fudos, *Depth-Fighting Aware Methods for Multi-fragment Rendering*, IEEE TVCG, vol. 19, no. 6, pp. 967-977, June, 2013.

Extracting coplanar fragments

Z-fighting phenomenon

more

- two or more primitives have same depth
- intersecting or overlapping surfaces
- multifragment rasterization is even more susceptible



Our contribution[†]

- ① [robust] & [approximate] methods
- ② [geometry culling] mechanism for multi-pass rendering

[†]A. Vasilakis & I. Fudos, *Depth-Fighting Aware Methods for Multi-fragment Rendering*, IEEE TVCG, vol. 19, no. 6, pp. 967-977, June, 2013.

Correcting Raster-based Pipelines

Adapting depth peeling methods based on

- ① primitive identifiers

Correcting Raster-based Pipelines

Adapting depth peeling methods based on

- ① primitive identifiers
- ② buffer-based solutions

Correcting Raster-based Pipelines

Adapting depth peeling methods based on

- ① primitive identifiers
- ② buffer-based solutions
- MSAA - Tessellation - Instancing

Correcting Raster-based Pipelines

Adapting depth peeling methods based on

- ① primitive identifiers
- ② buffer-based solutions
- MSAA - Tessellation - Instancing

Robustness ratio

- ① [robust]: Ma: low [\uparrow] - P: slow [\downarrow]

Correcting Raster-based Pipelines

Adapting depth peeling methods based on

- ① primitive identifiers
- ② buffer-based solutions
- MSAA - Tessellation - Instancing

Robustness ratio

- ① [robust]: Ma: low [\uparrow] - P: slow [\downarrow]
- ② [approximate]: Ma: high [\downarrow] - P: efficient [\uparrow]

Robust Algorithms (1)

Extending F2B¹,DUAL² → F2B-2P,DUAL-2P

example

- base methods extract only **one** coplanar fragment

Robust Algorithms (1)

Extending F2B¹,DUAL² → F2B-2P,DUAL-2P

example

- base methods extract only **one** coplanar fragment
- extracts 2 fragments/iteration **Ma**: constant (low)

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08]

Robust Algorithms (1)

Extending F2B¹,DUAL² → F2B-2P,DUAL-2P

example

- base methods extract only **one** coplanar fragment
- extracts 2 fragments/iteration **Ma**: constant (low)
- **[Idea]**: extra accumulation rendering pass
 - primitive ID (OpenGL: `gl_PrimitiveID`)
 - store min/max IDs of the remaining non-peeled fragments:
 $id_{\min}^1 < id_{\min}^2 < \dots < id_{\max}^2 < id_{\max}^1$

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08]

Robust Algorithms (1)

Extending F2B¹,DUAL² → F2B-2P,DUAL-2P

example

- base methods extract only **one** coplanar fragment
- extracts 2 fragments/iteration **Ma**: constant (low)
- **[Idea]**: extra accumulation rendering pass
 - primitive ID (OpenGL: `gl_PrimitiveID`)
 - store min/max IDs of the remaining non-peeled fragments:
 $id_{\min}^1 < id_{\min}^2 < \dots < id_{\max}^2 < id_{\max}^1$
- subsequent rendering pass:
 - ➊ extract fragment information using captured IDs

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08]

Robust Algorithms (1)

Extending F2B¹,DUAL² → F2B-2P,DUAL-2P

example

- base methods extract only **one** coplanar fragment
- extracts 2 fragments/iteration **Ma**: constant (low)
- **[Idea]**: extra accumulation rendering pass
 - primitive ID (OpenGL: `gl_PrimitiveID`)
 - store min/max IDs of the remaining non-peeled fragments:
 $id_{\min}^1 < id_{\min}^2 < \dots < id_{\max}^2 < id_{\max}^1$
- subsequent rendering pass:
 - ① extract fragment information using captured IDs
 - ② move or not to next depth layer (fragment coplanarity counter)

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08]

Robust Algorithms (1)

example

Extending F2B¹,DUAL² → F2B-2P,DUAL-2P

- base methods extract only **one** coplanar fragment
- extracts 2 fragments/iteration **Ma**: constant (low)
- **[Idea]**: extra accumulation rendering pass
 - primitive ID (OpenGL: `gl_PrimitiveID`)
 - store min/max IDs of the remaining non-peeled fragments:
 $id_{\min}^1 < id_{\min}^2 < \dots < id_{\max}^2 < id_{\max}^1$
- subsequent rendering pass:
 - ➊ extract fragment information using captured IDs
 - ➋ move or not to next depth layer (fragment coplanarity counter)

Extending F2B-2P → F2B-3P

- additional rendering pass: (double speed depth, early-z culling)

Robust Algorithms (1)

example

Extending F2B¹,DUAL² → F2B-2P,DUAL-2P

- base methods extract only **one** coplanar fragment
- extracts 2 fragments/iteration **Ma**: constant (low)
- **[Idea]**: extra accumulation rendering pass
 - primitive ID (OpenGL: `gl_PrimitiveID`)
 - store min/max IDs of the remaining non-peeled fragments:
 $id_{\min}^1 < id_{\min}^2 < \dots < id_{\max}^2 < id_{\max}^1$
- subsequent rendering pass:
 - ① extract fragment information using captured IDs
 - ② move or not to next depth layer (fragment coplanarity counter)

Extending F2B-2P → F2B-3P

- additional rendering pass: (double speed depth, early-z culling)
- **P**: \uparrow **Ma**: $-[-]$

Robust Algorithms (2)

Combining F2B¹,DUAL² with AB_{LL}³ → F2B-LL,DUAL-LL

example

- handle fragment coplanarity of arbitrary length per pixel

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Yang et al, CGF'10]

Robust Algorithms (2)

Combining F2B¹,DUAL² with AB_{LL}³ → F2B-LL,DUAL-LL

example

- handle fragment coplanarity of arbitrary length per pixel
- rendering workflow (2 passes/depth layer)
 - ➊ double speed depth pass
 - ➋ fragment linked lists at the current depth layer

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Yang et al, CGF'10]

Robust Algorithms (2)

Combining F2B¹,DUAL² with AB_{LL}³ → F2B-LL,DUAL-LL

example

- handle fragment coplanarity of arbitrary length per pixel
- rendering workflow (2 passes/depth layer)
 - ① double speed depth pass
 - ② fragment linked lists at the current depth layer

Limitations

- P: [↓]
 - ① multi-pass rendering
 - ② fragment storing (reduced congestion)
- G: only modern cards - [↓]

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Yang et al, CGF'10]

Robust Algorithms (3)

Combining Uniform Buckets¹ with AB_{LL} ² \rightarrow BUN-LL

example

- single-pass nature

¹ [Liu, HPG'09], ² [Yang et al, CGF'10]

Robust Algorithms (3)

Combining Uniform Buckets¹ with AB_{LL}² → BUN-LL

example

- single-pass nature
- uniformly split of the depth range
 - render bounding box
 - max 5 consecutive subintervals

¹ [Liu, HPG'09], ² [Yang et al, CGF'10]

Robust Algorithms (3)

Combining Uniform Buckets¹ with AB_{LL} ² \rightarrow BUN-LL

example

- single-pass nature
- uniformly split of the depth range
 - render bounding box
 - max 5 consecutive subintervals
- assign a linked list to each subdivision

¹ [Liu, HPG'09], ² [Yang et al, CGF'10]

Robust Algorithms (3)

Combining Uniform Buckets¹ with AB_{LL}² → BUN-LL

example

- single-pass nature
- uniformly split of the depth range
 - render bounding box
 - max 5 consecutive subintervals
- assign a linked list to each subdivision

Limitations

- P:
 - always faster than AB_{LL} – [↑]
 - uniform depth distribution - [↓]

¹ [Liu, HPG'09], ² [Yang et al, CGF'10]

Robust Algorithms (3)

Combining Uniform Buckets¹ with AB_{LL}² → BUN-LL

example

- single-pass nature
- uniformly split of the depth range
 - render bounding box
 - max 5 consecutive subintervals
- assign a linked list to each subdivision

Limitations

- **P:**
 - always faster than AB_{LL} – [↑]
 - uniform depth distribution - [↓]
- **Ma:** (extra memory) - [↓], **Mc:** [↓]

¹ [Liu, HPG'09], ² [Yang et al, CGF'10]

Robust Algorithms (3)

Combining Uniform Buckets¹ with AB_{LL}² → BUN-LL

example

- single-pass nature
- uniformly split of the depth range
 - render bounding box
 - max 5 consecutive subintervals
- assign a linked list to each subdivision

Limitations

- **P:**
 - always faster than AB_{LL} – [↑]
 - uniform depth distribution - [↓]
- **Ma:** (extra memory) - [↓], **Mc:** [↓]
- **G:** only modern cards - [↓]

¹ [Liu, HPG'09], ² [Yang et al, CGF'10]

Approximate Algorithms

Combining F2B¹,DUAL² with fixed-size arrays → F2B-B,DUAL-B

example

- G: modern - $[AB_{FP}]^3 \rightarrow$ F2B-FP,DUAL-FP

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, I3D'10], ⁴ [Baboi et al, I3D'07]

Approximate Algorithms

Combining F2B¹,DUAL² with fixed-size arrays → F2B-B,DUAL-B

example

- G: modern - $[AB_{FP}]^3$ → F2B-FP,DUAL-FP
 - bounded-length vectors per pixel

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, I3D'10], ⁴ [Baboi et al, I3D'07]

Approximate Algorithms

Combining F2B¹,DUAL² with fixed-size arrays → F2B-B,DUAL-B

example

- G: modern - $[AB_{FP}]^3$ → F2B-FP,DUAL-FP
 - bounded-length vectors per pixel
 - A: 100% - $[\uparrow]$
 - ① max{coplanar fragments/depth layer} known

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, I3D'10], ⁴ [Baboi et al, I3D'07]

Approximate Algorithms

Combining F2B¹,DUAL² with fixed-size arrays → F2B-B,DUAL-B

example

- G: modern - $[AB_{FP}]^3$ → F2B-FP,DUAL-FP
 - bounded-length vectors per pixel
 - A: 100% - $[\uparrow]$
 - ① max{coplanar fragments/depth layer} known
 - ② no memory overflow

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, I3D'10], ⁴ [Baboi et al, I3D'07]

Approximate Algorithms

Combining F2B¹,DUAL² with fixed-size arrays → F2B-B,DUAL-B

example

- G: modern - [AB_{FP}]³ → F2B-FP,DUAL-FP
 - bounded-length vectors per pixel
 - A: 100% - [\uparrow]
 - ① max{coplanar fragments/depth layer} known
 - ② no memory overflow
- G: commodity - [KB]⁴ → F2B-KB,DUAL-KB

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, I3D'10], ⁴ [Baboi et al, I3D'07]

Approximate Algorithms

Combining F2B¹,DUAL² with fixed-size arrays → F2B-B,DUAL-B

example

- G: modern - $[AB_{FP}]^3$ → F2B-FP,DUAL-FP
 - bounded-length vectors per pixel
 - A: 100% - $[\uparrow]$
 - ① max{coplanar fragments/depth layer} known
 - ② no memory overflow
- G: commodity - $[KB]^4$ → F2B-KB,DUAL-KB
 - 8 coplanar fragments/layer

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, I3D'10], ⁴ [Baboi et al, I3D'07]

Approximate Algorithms

Combining F2B¹,DUAL² with fixed-size arrays → F2B-B,DUAL-B

example

- G: modern - [AB_{FP}]³ → F2B-FP,DUAL-FP
 - bounded-length vectors per pixel
 - A: 100% - [\uparrow]
 - ① max{coplanar fragments/depth layer} known
 - ② no memory overflow
- G: commodity - [KB]⁴ → F2B-KB,DUAL-KB
 - 8 coplanar fragments/layer
 - [data packing]: 32 coplanar fragments/layer

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, I3D'10], ⁴ [Baboi et al, I3D'07]

Approximate Algorithms

Combining F2B¹,DUAL² with fixed-size arrays → F2B-B,DUAL-B

example

- G: modern - $[AB_{FP}]^3$ → F2B-FP,DUAL-FP
 - bounded-length vectors per pixel
 - A: 100% - $[\uparrow]$
 - ① max{coplanar fragments/depth layer} known
 - ② no memory overflow
- G: commodity - $[KB]^4$ → F2B-KB,DUAL-KB
 - 8 coplanar fragments/layer
 - [data packing]: 32 coplanar fragments/layer
 - no sorting needed: [RMW hazard-free]

¹ [Everitt, Nvidia'01], ² [Bavoil et al, Nvidia'08], ³ [Liu et al, I3D'10], ⁴ [Baboi et al, I3D'07]

Optimizing multi-pass rendering of multiple objects

Occlusion culling mechanism

- geometry is not rendered when is hidden by objects closer to camera

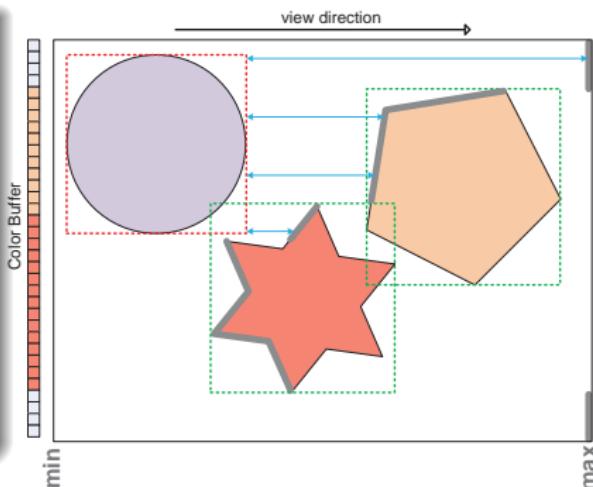
Optimizing multi-pass rendering of multiple objects

Occlusion culling mechanism

- geometry is not rendered when is hidden by objects closer to camera

Avoid rendering *completely-peeled* objects

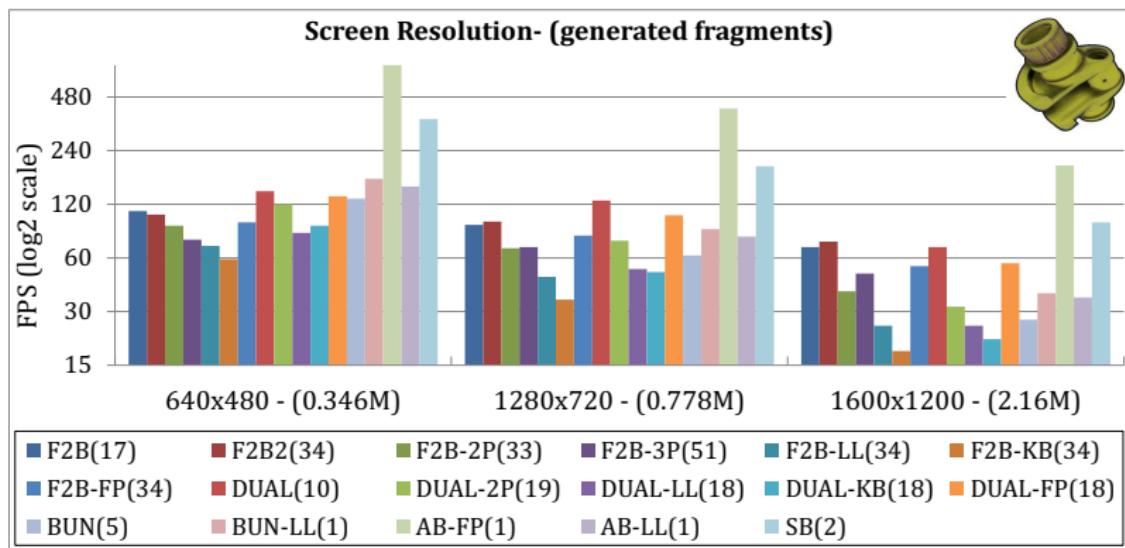
- [Goal]: rendering load reduction of the following passes
- if objects bounding box is behind current depth layer then [cull]
- hardware occlusion queries
- reuse query results from previous iterations



Performance Analysis (1)

Impact of Screen Resolution

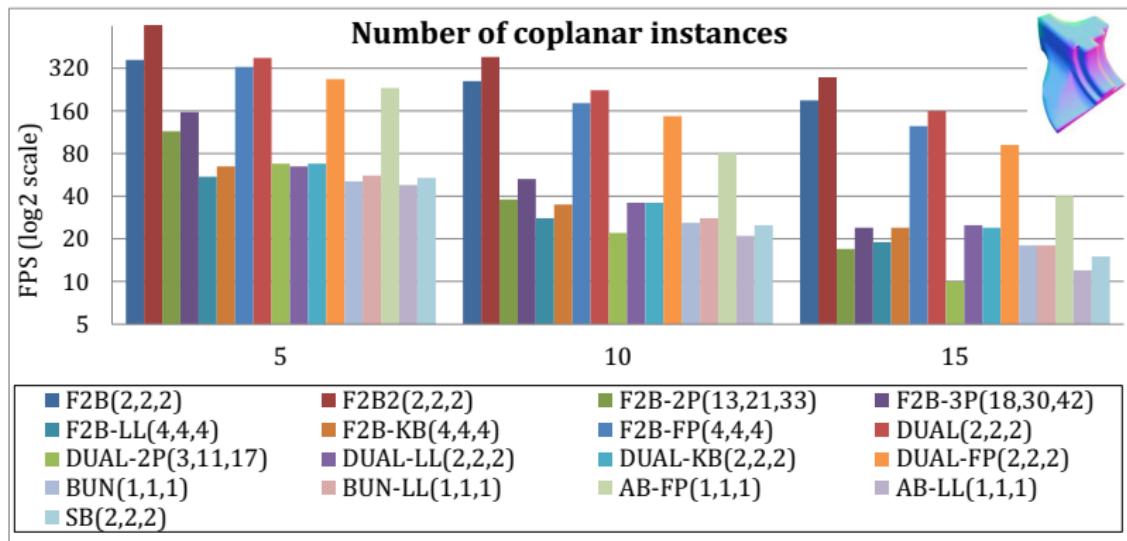
- [Crank]: 10K triangles, 17 depth layers, no coplanarity



Performance Analysis (2)

Impact of Coplanarity

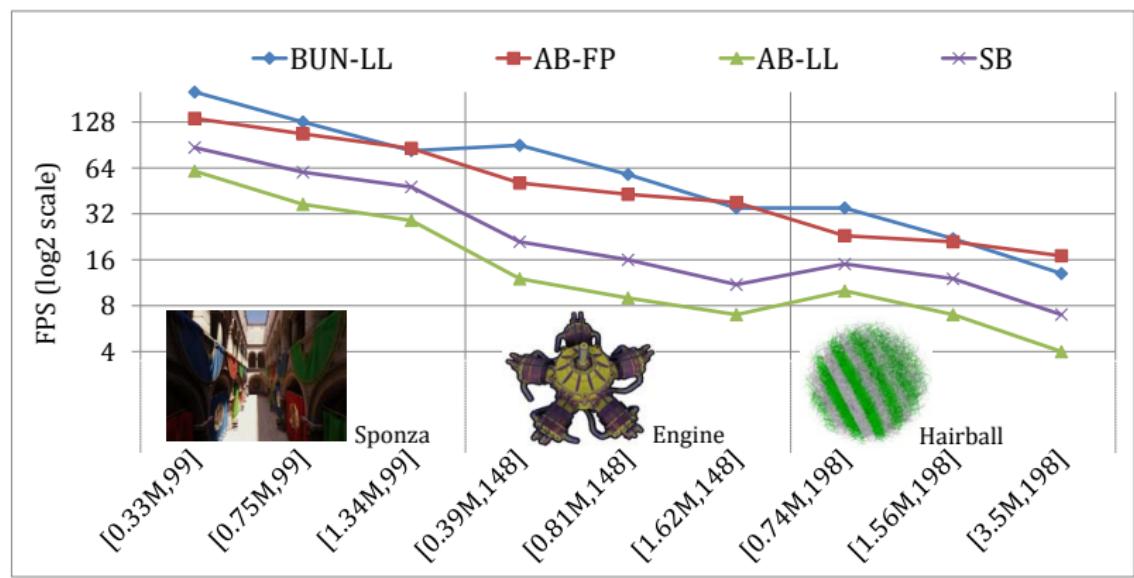
- [Fandisk]: 2K triangles, 2 depth layers, fragments/layer=instances



Performance Analysis (3)

Impact of High Depth Complexity

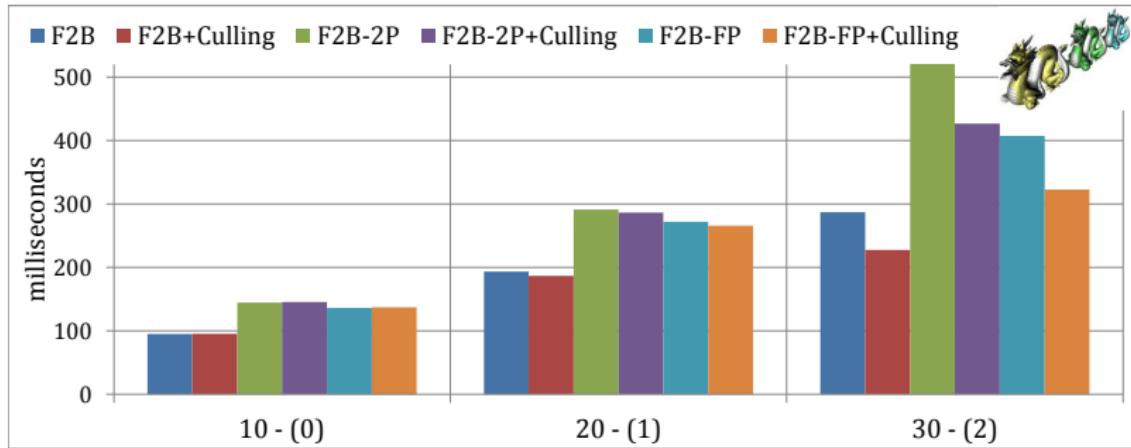
- [Sponza] (279K) [Engine] (203K) [Hairball] (2.85M) triangles



Performance Analysis (4)

Impact of Culling

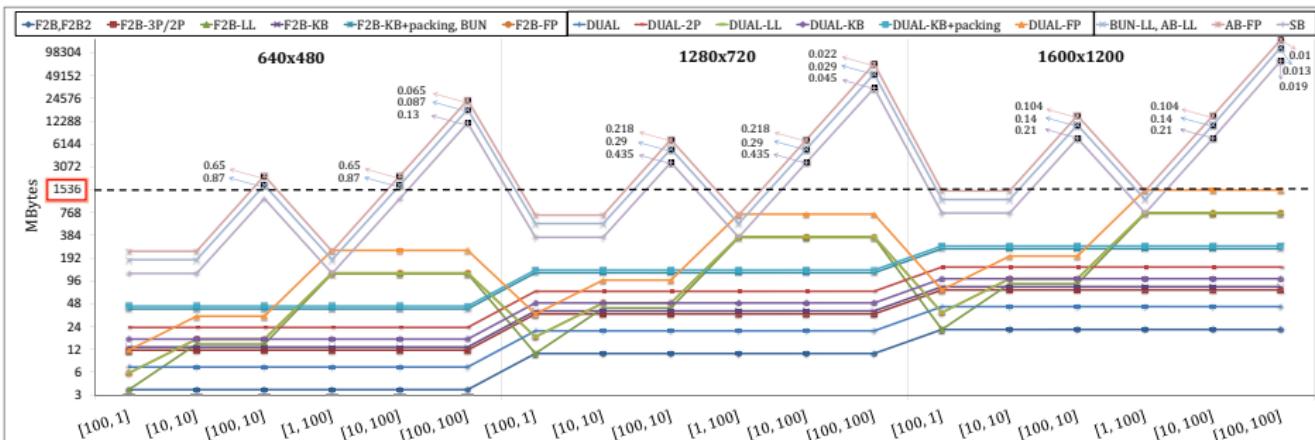
- [Dragon]: 870K triangles, 10 depth layers



Memory Allocation Analysis

Impact of # Generated Fragments

- [Robustness ratio] ???



Discussion

G: modern hardware

- [Low Memory]: Winner(FAB)
- [Medium Memory]:
 - [Low depth complexity]: Winner(SB)
 - [High depth complexity]: Winner(BUN-LL)
- [High Memory]:
 - [Low coplanarity]: Winner(F2B-FAB, DUAL-FAB)
 - [High coplanarity]: Winner(F2B-LL, DUAL-LL)

Discussion

G: modern hardware

- [Low Memory]: Winner(FAB)
- [Medium Memory]:
 - [Low depth complexity]: Winner(SB)
 - [High depth complexity]: Winner(BUN-LL)
- [High Memory]:
 - [Low coplanarity]: Winner(F2B-FAB, DUAL-FAB)
 - [High coplanarity]: Winner(F2B-LL, DUAL-LL)

G: commodity hardware

- [Low coplanarity]: Winner(F2B-3P, DUAL-2P)
- [High coplanarity]: Winner(F2B-KB, DUAL-KB)

Discussion

G: modern hardware

- [Low Memory]: Winner(FAB)
- [Medium Memory]:
 - [Low depth complexity]: Winner(SB)
 - [High depth complexity]: Winner(BUN-LL)
- [High Memory]:
 - [Low coplanarity]: Winner(F2B-FAB, DUAL-FAB)
 - [High coplanarity]: Winner(F2B-LL, DUAL-LL)

G: commodity hardware

- [Low coplanarity]: Winner(F2B-3P, DUAL-2P)
- [High coplanarity]: Winner(F2B-KB, DUAL-KB)

F2B versus DUAL extensions

- [Tessellation]: low/high - [Resolution]: low/high

Outline

- 1 Background and Preliminaries
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 S-buffer: Sparsity-aware multi-fragment rendering
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 **k^+ -buffer: Fragment synchronized k -buffer**
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

- capture the k -closest fragments

¹ [Bavoil et al, I3D'07], ² [Mayers et al, SIGGRAPH'07], ³ [Maule et al, I3D'13], ⁴ [Yu et al, I3D'12]

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

- capture the k -closest fragments
- reduce memory & sorting costs

¹ [Bavoil et al, I3D'07], ² [Mayers et al, SIGGRAPH'07], ³ [Maule et al, I3D'13], ⁴ [Yu et al, I3D'12]

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

- capture the k -closest fragments
- reduce memory & sorting costs
- suffer from
 - ① RMW hazards¹
 - ② geometry pre-sorting^{1,2}

¹ [Bavoil et al, I3D'07], ² [Mayers et al, SIGGRAPH'07], ³ [Maule et al, I3D'13], ⁴ [Yu et al, I3D'12]

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

- capture the k -closest fragments
- reduce memory & sorting costs
- suffer from
 - ① RMW hazards¹
 - ② geometry pre-sorting^{1,2}
 - ③ depth precision conversion³

¹ [Bavoil et al, I3D'07], ² [Mayers et al, SIGGRAPH'07], ³ [Maule et al, I3D'13], ⁴ [Yu et al, I3D'12]

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

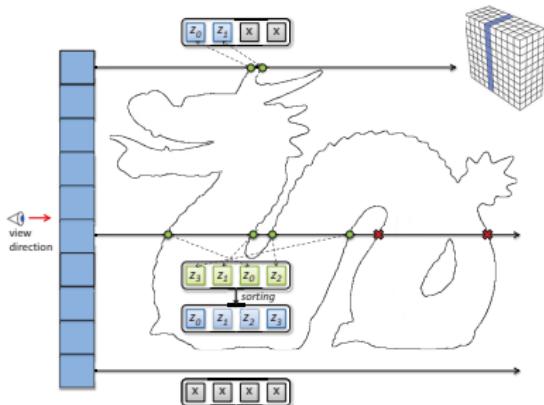
- capture the k -closest fragments
- reduce memory & sorting costs
- suffer from
 - ① RMW hazards¹
 - ② geometry pre-sorting^{1,2}
 - ③ depth precision conversion³
 - ④ unbounded memory⁴

¹ [Bavoil et al, I3D'07], ² [Mayers et al, SIGGRAPH'07], ³ [Maule et al, I3D'13], ⁴ [Yu et al, I3D'12]

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

- capture the k -closest fragments
- reduce memory & sorting costs
- suffer from
 - RMW hazards¹
 - geometry pre-sorting^{1,2}
 - depth precision conversion³
 - unbounded memory⁴

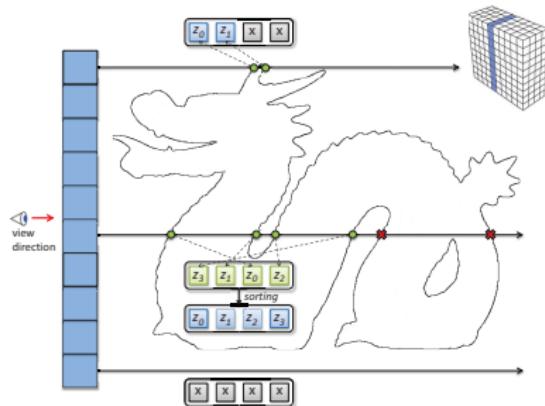


¹ [Bavoil et al, I3D'07], ² [Mayers et al, SIGGRAPH'07], ³ [Maule et al, I3D'13], ⁴ [Yu et al, I3D'12]

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

- capture the k -closest fragments
- reduce memory & sorting costs
- suffer from
 - RMW hazards¹
 - geometry pre-sorting^{1,2}
 - depth precision conversion³
 - unbounded memory⁴



Our contribution - k^+ -buffer[†]

- overcomes all limitations of existing k -buffer alternatives

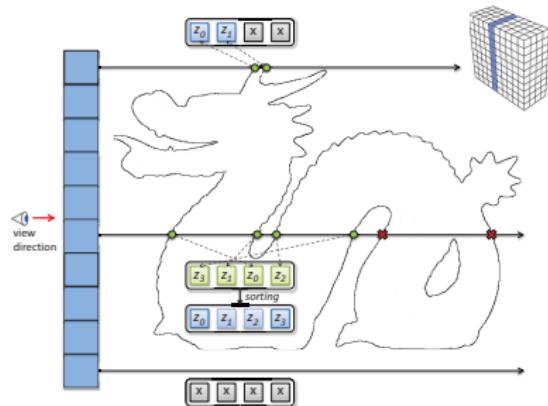
[†]A. Vasilakis & I. Fudos, k^+ -buffer: Fragment synchronized k -buffer,

Symposium on Interactive 3D Graphics and Games (I3D 14), San Francisco, USA, March, 2014.

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

- capture the k -closest fragments
- reduce memory & sorting costs
- suffer from
 - RMW hazards¹
 - geometry pre-sorting^{1,2}
 - depth precision conversion³
 - unbounded memory⁴



Our contribution - k^+ -buffer[†]

- overcomes all limitations of existing k -buffer alternatives
- memory-friendly variation (extra pass needed)

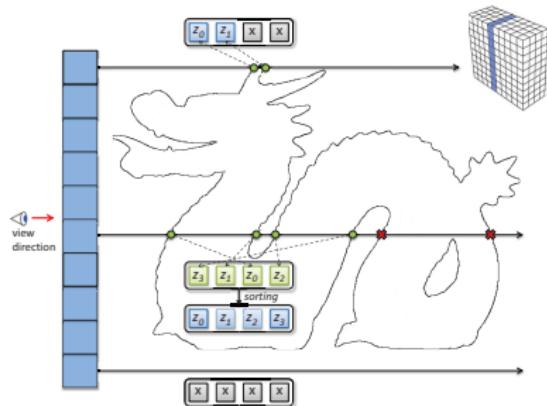
[†]A. Vasilakis & I. Fudos, k^+ -buffer: Fragment synchronized k -buffer,

Symposium on Interactive 3D Graphics and Games (I3D 14), San Francisco, USA, March, 2014.

Revisiting k -buffer

k -buffer¹ & variants^{2,3,4}

- capture the k -closest fragments
- reduce memory & sorting costs
- suffer from
 - RMW hazards¹
 - geometry pre-sorting^{1,2}
 - depth precision conversion³
 - unbounded memory⁴



Our contribution - k^+ -buffer[†]

- overcomes all limitations of existing k -buffer alternatives
- memory-friendly variation (extra pass needed)
- supports Z-buffer and A-buffer functionalities

[†]A. Vasilakis & I. Fudos, k^+ -buffer: Fragment synchronized k -buffer,

Symposium on Interactive 3D Graphics and Games (I3D 14), San Francisco, USA, March, 2014.

k^+ -buffer Pipeline

Store & Sort solution:

- ① captures fragments in an unsorted sequence

k^+ -buffer Pipeline

Store & Sort solution:

- ① captures fragments in an unsorted sequence
- ② reorders stored fragments by their depth

Store Pass:

example

- ① spin-lock strategy ([binary semaphores]¹ or [pixel sync]²)
 - avoids 32-bit atomic operations

¹ [Crassin, blog'10], ² [Salvi, SIGGRAPH'13]

k^+ -buffer Pipeline

Store & Sort solution:

- ① captures fragments in an unsorted sequence
- ② reorders stored fragments by their depth

Store Pass:

example

- ① spin-lock strategy ([binary semaphores]¹ or [pixel sync]²)
 - avoids 32-bit atomic operations
- ② two bounded array-based structures:
 - if $k < 16$: [max-array]-($K^+ B$ -Array), else [max-heap]-($K^+ B$ -Heap)

¹ [Crassin, blog'10], ² [Salvi, SIGGRAPH'13]

k^+ -buffer Pipeline

Store & Sort solution:

- ① captures fragments in an unsorted sequence
- ② reorders stored fragments by their depth

Store Pass:

example

- ① spin-lock strategy ([binary semaphores]¹ or [pixel sync]²)
 - avoids 32-bit atomic operations
- ② two bounded array-based structures:
 - if $k < 16$: [max-array]-($K^+ B$ -Array), else [max-heap]-($K^+ B$ -Heap)
 - early-fragment culling

¹ [Crassin, blog'10], ² [Salvi, SIGGRAPH'13]

k^+ -buffer Pipeline

Store & Sort solution:

- ① captures fragments in an unsorted sequence
- ② reorders stored fragments by their depth

Store Pass:

example

- ① spin-lock strategy ([binary semaphores]¹ or [pixel sync]²)
 - avoids 32-bit atomic operations
- ② two bounded array-based structures:
 - if $k < 16$: [max-array]-($K^+ B$ -Array), else [max-heap]-($K^+ B$ -Heap)
 - early-fragment culling

Sorting Pass:

- if $k < 16$: [insertion-sort], else [shell-sort]

¹ [Crassin, blog'10], ² [Salvi, SIGGRAPH'13]

Extending k^+ -buffer pipeline

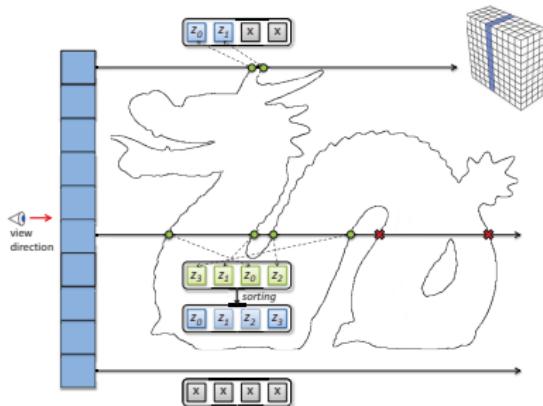
Precise memory allocation

- k is the same for all pixels

Extending k^+ -buffer pipeline

Precise memory allocation

- k is the same for all pixels
- [Idea]: S-buffer¹ - ($K^+ B$ -SB)
 - counting \mapsto [hybrid scheme]
 - [condition] if k reaches n stop
 - extra pass & shared memory

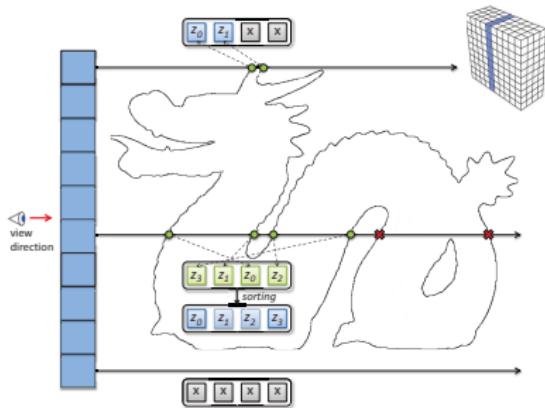


¹ [Vasilakis et al, EG'12], ² [Liu et al, I3D'10]

Extending k^+ -buffer pipeline

Precise memory allocation

- k is the same for all pixels
- [Idea]: S-buffer¹ - ($K^+ B$ -SB)
 - counting \mapsto [hybrid scheme]
 - [condition] if k reaches n stop
 - extra pass & shared memory



Unified framework

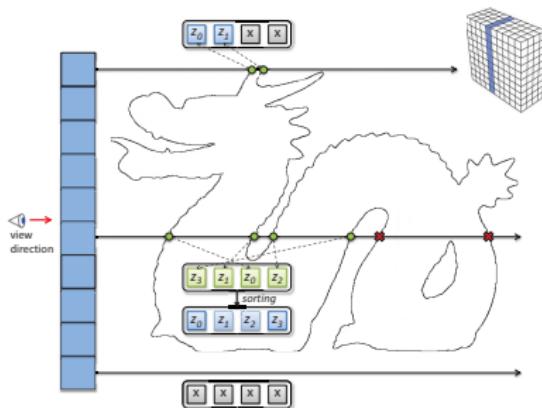
- adjust the value of k :
 - $k = 1$: [Z-buffer] behavior

¹ [Vasilakis et al, EG'12], ² [Liu et al, I3D'10]

Extending k^+ -buffer pipeline

Precise memory allocation

- k is the same for all pixels
- [Idea]: S-buffer¹ - ($K^+ B$ -SB)
 - counting \mapsto [hybrid scheme]
 - [condition] if k reaches n stop
 - extra pass & shared memory



Unified framework

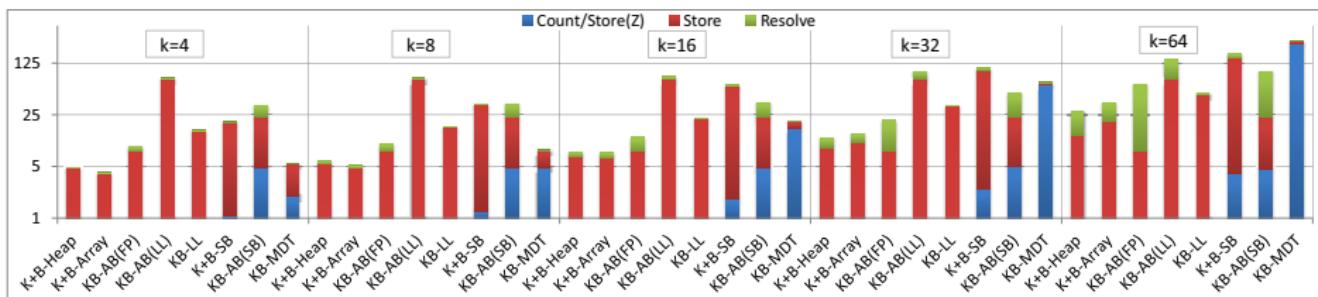
- adjust the value of k :
 - $k = 1$: [Z-buffer] behavior
 - $k = \max_p\{f(p)\}$: [A-buffer] behavior:
 - $K^+ B \mapsto AB_{FP}$ ²
 - $K^+ B$ -SB \mapsto S-buffer¹

¹ [Vasilakis et al, EG'12], ² [Liu et al, I3D'10]

Performance Analysis - k -buffer

Impact of k

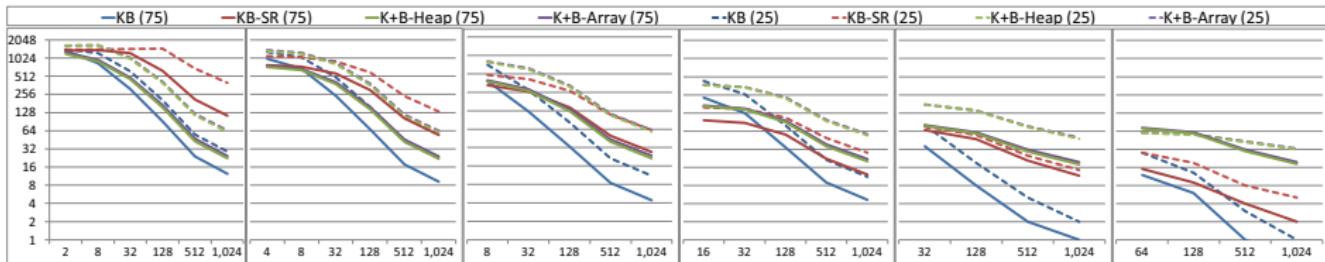
- [Scene]: $n = 128$, $k = \{4, \dots, 64\}$
- K^+B -Array [$k \downarrow$] - K^+B -Heap [$k \uparrow$]
- KB-MDT (two-pass method): future 64-bits?
- KB-AB_{FP}: storing [$k \uparrow$] - sorting [$k \downarrow$]



Performance Analysis - k -buffer

Impact of Sorting

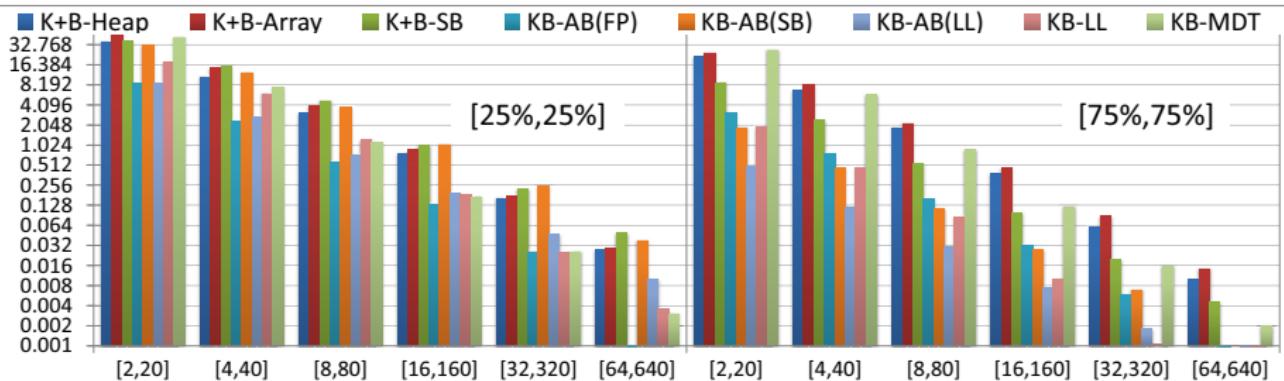
- [Scene]: $n = \{k, \dots, 1024\}$, $p_d = \{25\%, 75\%\}$, (in depth order)
- K^+B -Array $O(1) > K^+B$ -Heap $O(\log_2 k)$ - (linear behavior p_d)
- [$k \uparrow$] + multi-pass: $K^+B > KB-SR > KB$



Performance Analysis - k -buffer

Impact of Memory

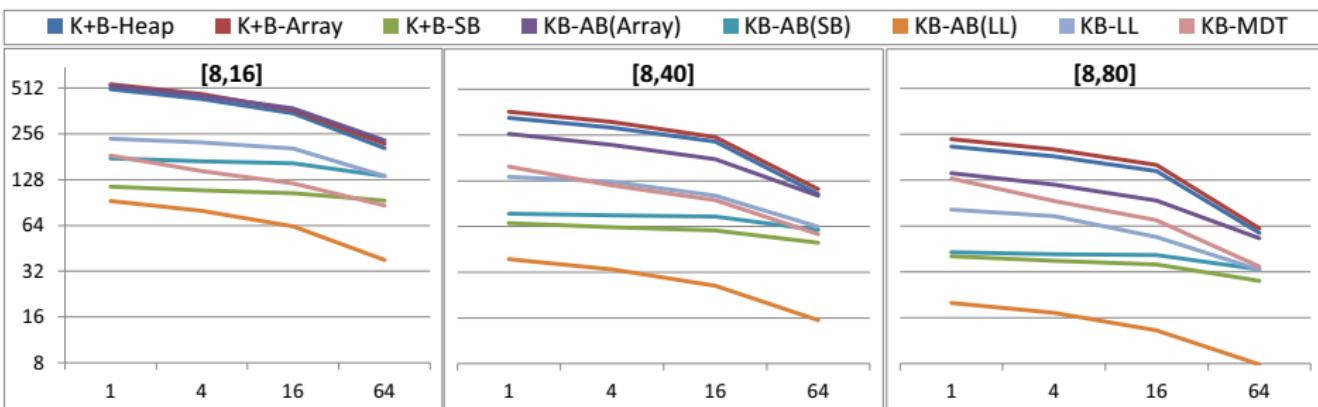
- [Scene]: $n = \{k, \dots, 10 \cdot k\}, [p_d, f_p]$ - speed: fps/MB
- $[p_d \downarrow, f_p \downarrow]$: $K^+B\text{-SB} > K^+B$
- $[p_d \uparrow, f_p \uparrow]$: $K^+B\text{-SB} < K^+B$
- A-buffer-based solutions fail (overflow) when $k = 64$



Performance Analysis - k -buffer

Impact of Tessellation

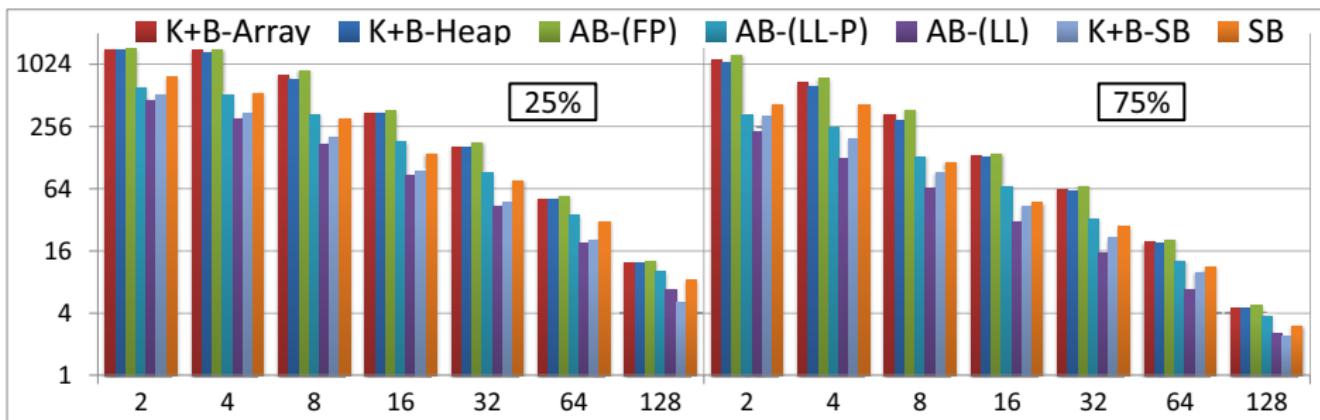
- [Scene]: level = $\{1, \dots, 64\}$, $k = 8, n = \{16, 40, 80\}$
- linear behavior of K^+B methods
- counting geometry pass is not-tessellation dependent



Performance Analysis - A-buffer

Impact of k

- [Scene]: $k = n, p_d = \{25\%, 75\%\}$
- $AB_{FP} > K^+B\text{-Array} > K^+B\text{-Heap} >$ rest methods (culling mechanism)
- $S\text{-buffer} > K^+B\text{-SB} > AB_{LL}$ (if condition at counting pass)



Memory Allocation Analysis

Comparison - bounded-buffers

- KB-AB_{FP} needs huge resources
- K⁺B methods require more storage (8-byte/pixel) than KB,KB-MHA
- [pixel sync] avoids semaphore allocation (4-byte)
- [data packing] employed: $\forall k > 1 : 4k > 2k + 2$
- K⁺B vs KB-SR: $\forall k > 2 : 3k > 2k + 2$

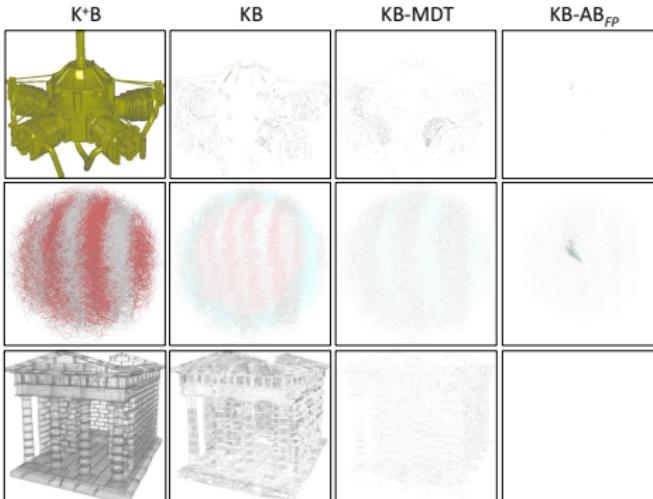
Comparison - unbounded-buffers

- K⁺B-SB requires:
$$\left. \begin{array}{l} \textcircled{1} \text{ [equal]}: (f(p) \leq k) \\ \textcircled{2} \text{ [less]}: (f(p) > k) \end{array} \right\} \text{KB-AB}_{LL}, \text{KB-LL}, \text{KB-AB}_{SB}$$
- A-buffer simulation: K⁺B = AB_{FP} & K⁺B-SB = S-buffer

Image Quality Analysis

Noticeable image differences

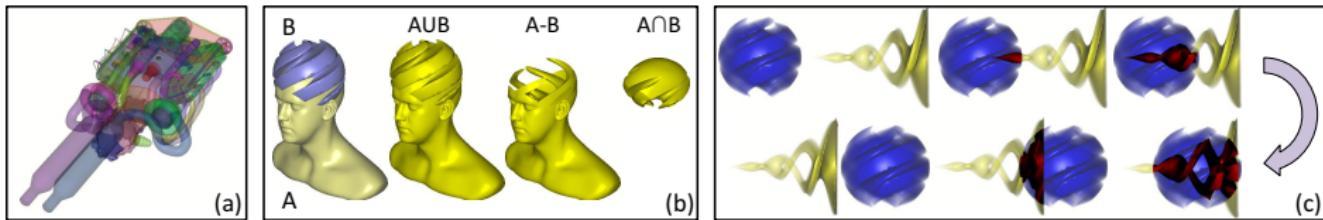
- (a) Z-buffer, (b) k-buffer, (c) A-buffer
- (left) KB: [RMW hazards], (middle) KB-MDT: [depth conversion],
(right) KB-AB_{FP}: [avoid overflow]



Conclusions

Bounded multi-fragment storage using k^+ -buffer

- alleviates prior k -buffer limitations and bottlenecks by exploiting fragment culling and pixel synchronization
- introduces an extension to avoid wasteful memory consumption.
- can also simulate the behavior of Z-buffer or A-buffer



Outline

- 1 Background and Preliminaries
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 S-buffer: Sparsity-aware multi-fragment rendering
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 k^+ -buffer: Fragment synchronized k -buffer
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Direct Rendering of Self-Trimmed Surfaces

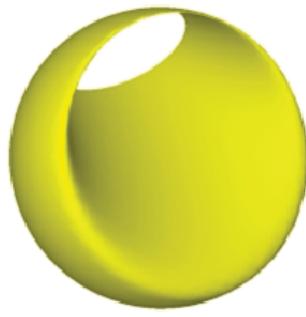
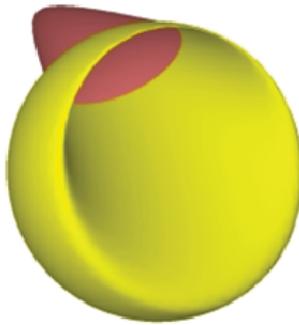
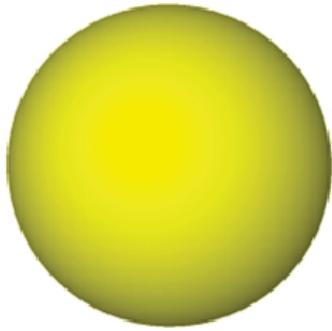
Rendering Self-Trimmed Surface (STS)

- provide interior definition (**I**) for self-crossing surface (SCS)

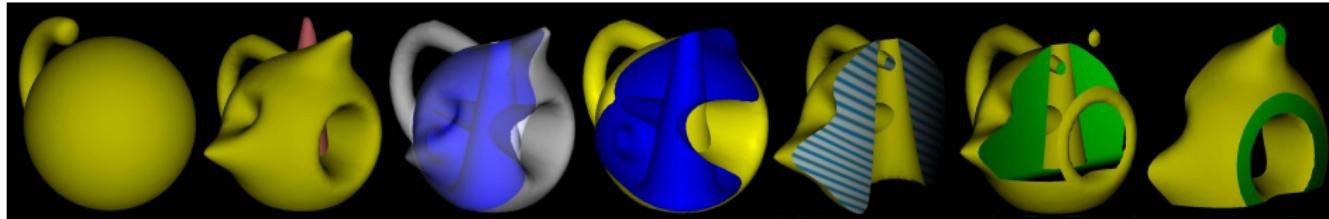
Direct Rendering of Self-Trimmed Surfaces

Rendering Self-Trimmed Surface (STS)

- provide interior definition (**I**) for self-crossing surface (SCS)
- render *trim* (boundary of **I**) of (a) [static] SCS & (b) [dynamic] SCS



Direct Rendering of Self-Trimmed Surfaces

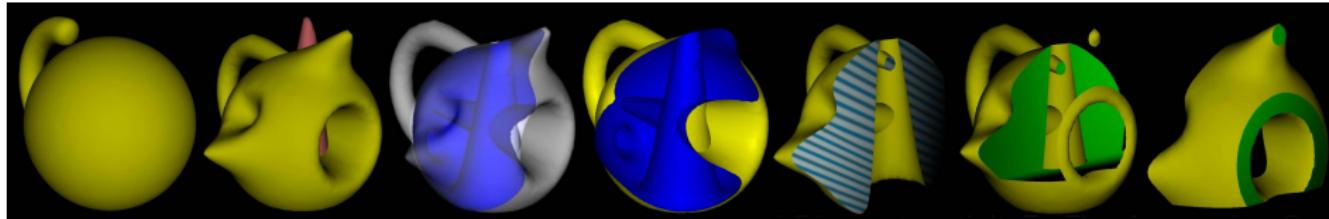


Our contribution[†]

- ① eliminate the need for computing self-crossing curves
- ② support intuitive animations & interactive editing results

[†] J. Rossignac & I. Fudos & **A. Vasilakis**, *Direct rendering of boolean combinations of self-trimmed surfaces*, Computer Aided Design, 45(2):288-300, 2013.

Direct Rendering of Self-Trimmed Surfaces

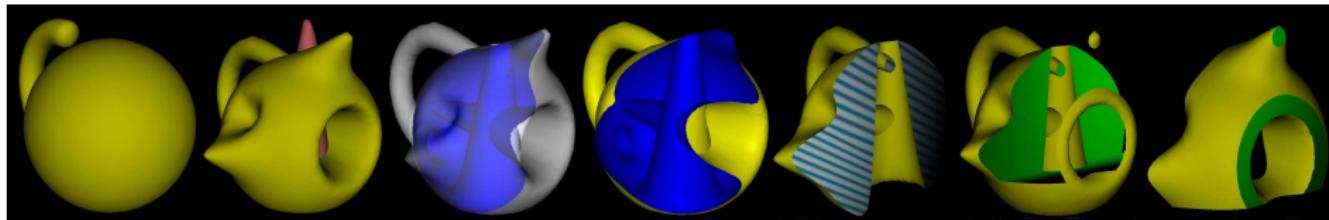


Our contribution[†]

- ① eliminate the need for computing self-crossing curves
- ② support intuitive animations & interactive editing results
- ③ provide real-time trimming & inspection tools for STS

[†] J. Rossignac & I. Fudos & A. Vasilakis, *Direct rendering of boolean combinations of self-trimmed surfaces*, Computer Aided Design, 45(2):288-300, 2013.

Direct Rendering of Self-Trimmed Surfaces



Our contribution[†]

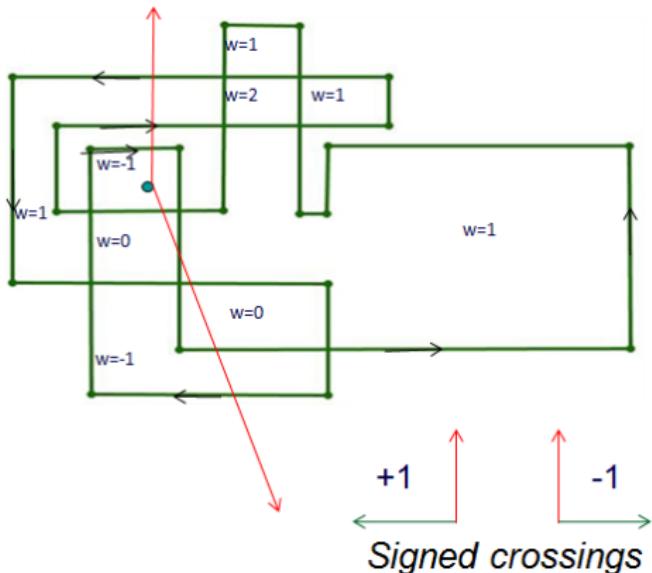
- ① eliminate the need for computing self-crossing curves
- ② support intuitive animations & interactive editing results
- ③ provide real-time trimming & inspection tools for STS
- ④ treat STS as valid primitives in CAD, including CSG

[†] J. Rossignac & I. Fudos & A. Vasilakis, *Direct rendering of boolean combinations of self-trimmed surfaces*, Computer Aided Design, 45(2):288-300, 2013.

Winding Number - 2D

Computation:

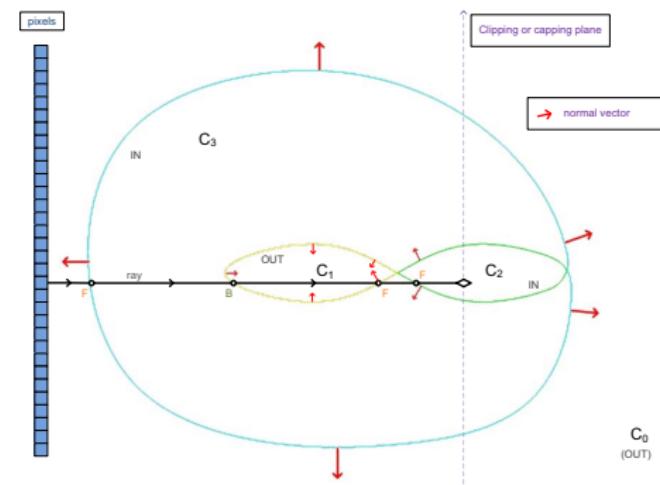
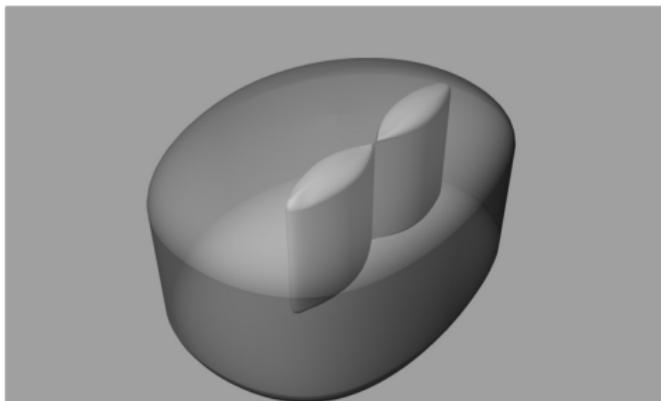
- count of signed crossings (+1 or -1) of ray to point outside
- independent of the choice ray



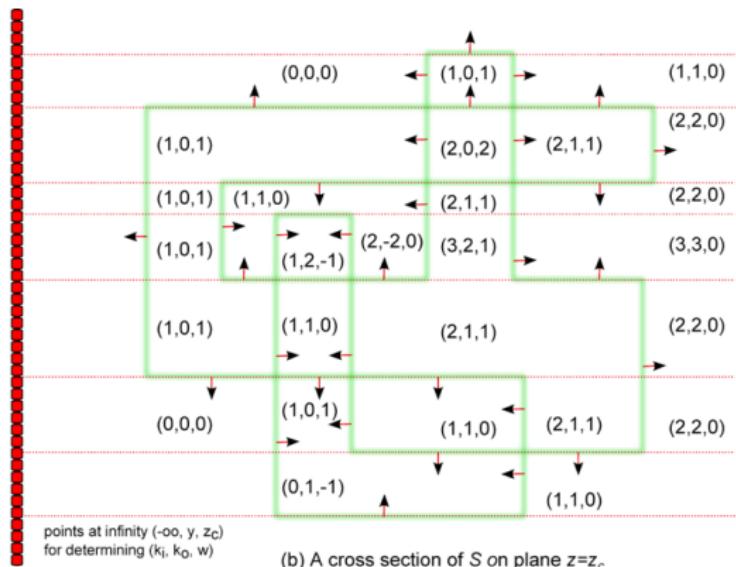
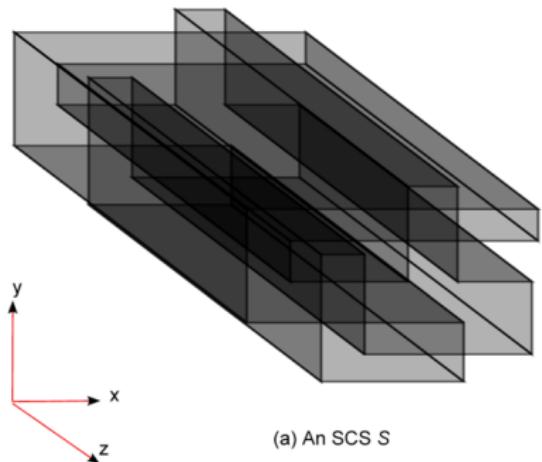
Winding Number - 3D

Computation:

- Surface S , Pixel P , Fragment f , point p
 - ① $k_i = \# P$ enters S ($f.\text{normal} [\uparrow]$)
 - ② $k_o = \# P$ exits S ($f.\text{normal} [\downarrow]$)
- $$w(p, S) = k_i - k_o$$



Winding Number - 3D Example



Static Rules

Given the winding number w :

Static Rules

Given the winding number w :

① [positive index]: $w > 0$

- not orientation invariant: fail if we reverse orientation

Static Rules

Given the winding number w :

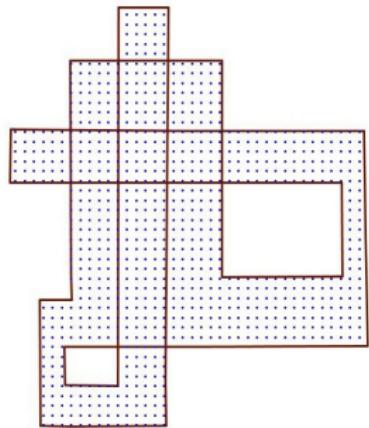
- ① [positive index]: $w > 0$
 - not orientation invariant: fail if we reverse orientation
- ② [parity] or [alternating component]: $w \% 2 = 1$ (or w is odd)
 - $T(S) = S$: does not trim the surface at all
 - produces non-manifolds solids

Static Rules

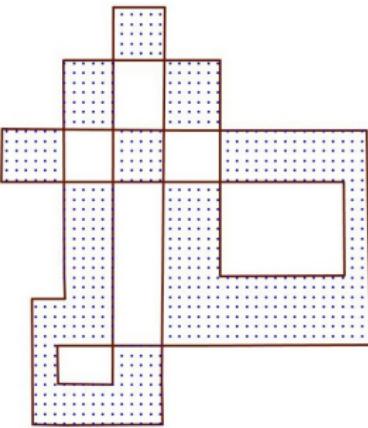
Given the winding number w :

- ① [positive index]: $w > 0$
 - not orientation invariant: fail if we reverse orientation
- ② [parity] or [alternating component]: $w \% 2 = 1$ (or w is odd)
 - $T(S) = S$: does not trim the surface at all
 - produces non-manifolds solids
- ③ [alternating border]: $\lceil \frac{w}{2} \rceil \% 2 = 1$
 - does not produce non-manifold solids
 - obtain complement of solid \mapsto add two surrounding boxes ($w + 2$)
 - reverse the surface orientation we obtain the complement of the trim

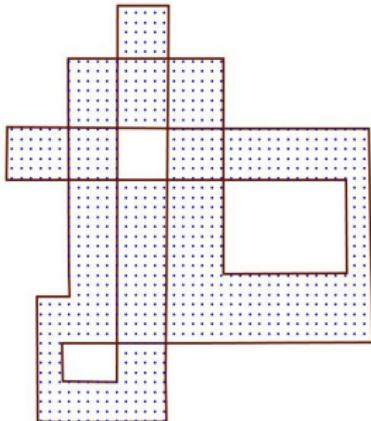
Static rules fail to capture well deformations...



(a)

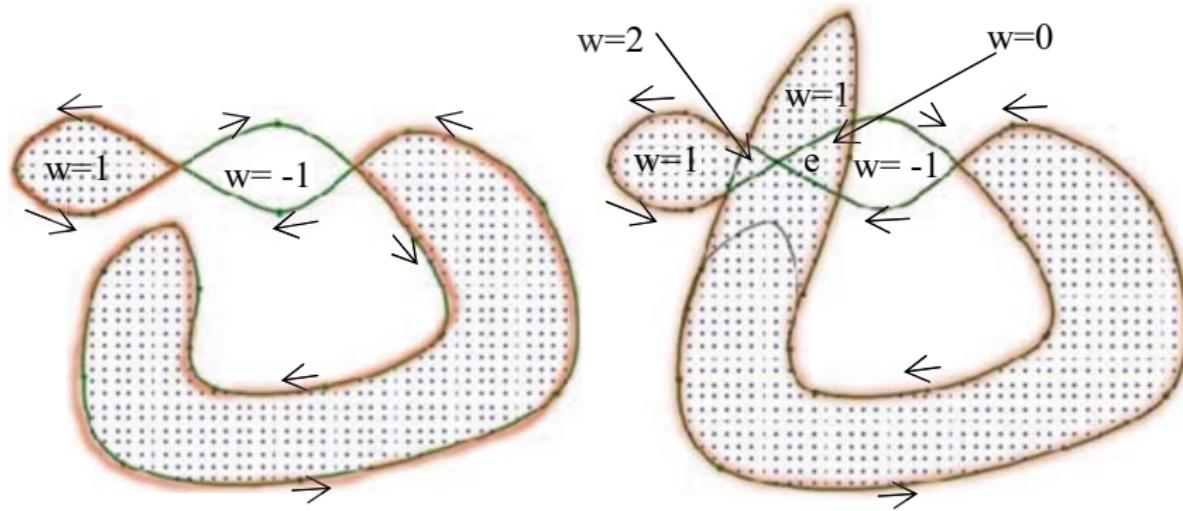


(b)

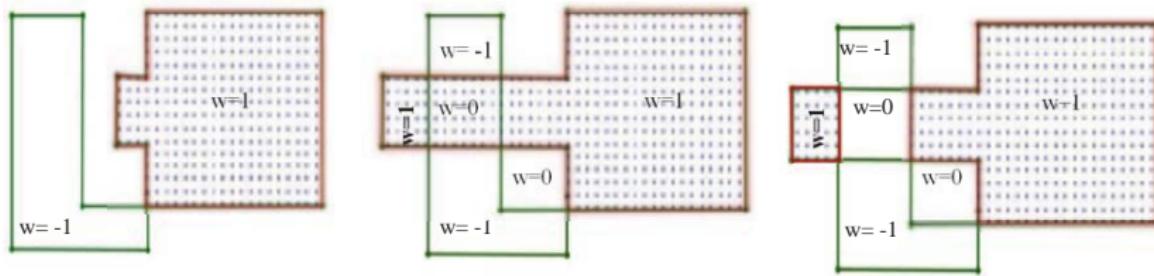


(c)

Static rules fail to capture well deformations...



Static rules fail to capture well deformations...



Dynamic Rules

Modeling Deformations for Dynamic Rules

- need to maintain history from reference S to deformed pose S'
- k disjoint deformations: a boundary surface may cross a point only once

Dynamic Rules

Modeling Deformations for Dynamic Rules

- need to maintain history from reference S to deformed pose S'
 - k disjoint deformations: a boundary surface may cross a point only once
- ① [constructive]: $i(p, S') = i(p, S)$ **op** ($w(p, S) \neq w(p, S')$)
- initial classification of S_0 : any static rule
 - design intent & has a constructive nature
 - ill-defined results: trim is not part of initial surface

Dynamic Rules

Modeling Deformations for Dynamic Rules

- need to maintain history from reference S to deformed pose S'
- k disjoint deformations: a boundary surface may cross a point only once

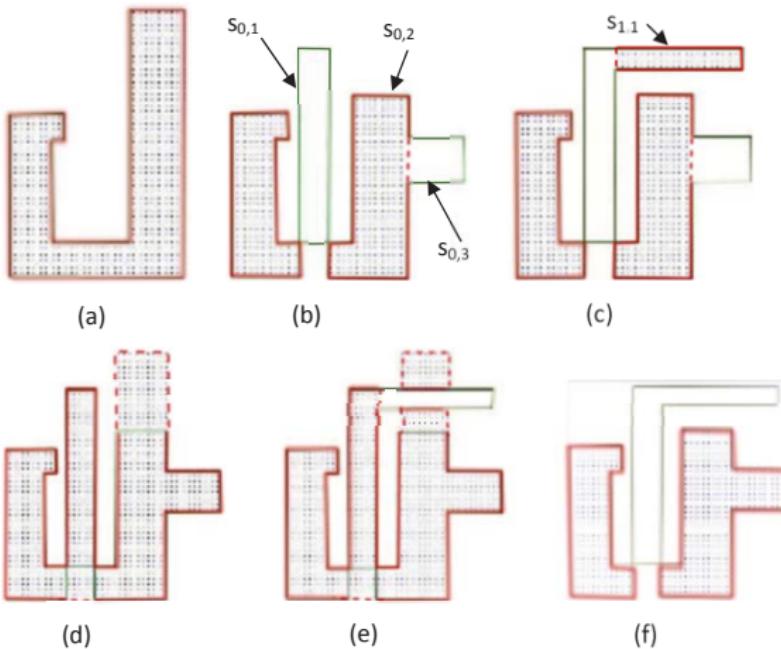
① [constructive]: $i(p, S') = i(p, S)$ **op** ($w(p, S) \neq w(p, S')$)

- initial classification of S_0 : any static rule
- design intent & has a constructive nature
- ill-defined results: trim is not part of initial surface

② [confluent deformation]: $i(p, S') = \begin{cases} i(p, S), w(p, S) = w(p, S') \\ \left\lceil \frac{w(p, S') - w(p, S)}{2} \right\rceil \% 2, \text{ otherwise} \end{cases}$

- prohibit deformations of the trimmed surface parts

Dynamic Rules - Example



Rendering & Trimming Algorithms for Static Rules

Pipeline

- ① [classification]: compute $p.\text{index} \longmapsto p.\text{class}$

Rendering & Trimming Algorithms for Static Rules

Pipeline

- ① [classification]: compute $p.\text{index} \mapsto p.\text{class}$
- ② [trimming]: find f with different in/out classifications

Rendering & Trimming Algorithms for Static Rules

Pipeline

- ① [classification]: compute $p.\text{index} \mapsto p.\text{class}$
- ② [trimming]: find f with different in/out classifications
- ③ [clipping]: start count/render after $C.\text{depth}$

Rendering & Trimming Algorithms for Static Rules

Pipeline

- ① [classification]: compute $p.\text{index} \mapsto p.\text{class}$
- ② [trimming]: find f with different in/out classifications
- ③ [clipping]: start count/render after $C.\text{depth}$

Multi-fragment Rendering

- ① multi-pass depth peeling
- ② buffer-based methods

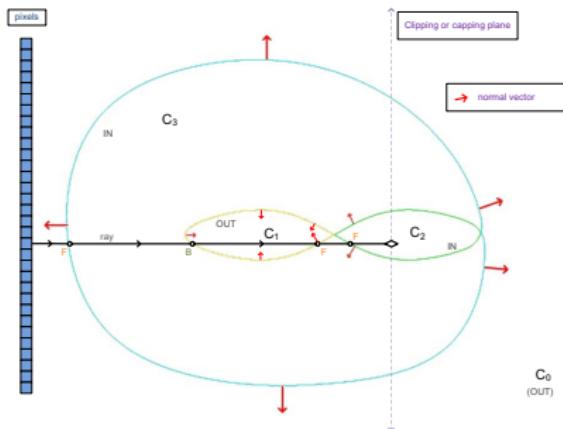
Rendering & Trimming Algorithms for Static Rules

Pipeline

- ① [classification]: compute $p.\text{index} \mapsto p.\text{class}$
- ② [trimming]: find f with different in/out classifications
- ③ [clipping]: start count/render after $C.\text{depth}$

Multi-fragment Rendering

- ① multi-pass depth peeling
- ② buffer-based methods
- ③ sort-independent methods¹



¹ [Meshkin, GDC'07]

Rendering & Trimming Algorithms for Dynamic Rules

Pipeline

- ➊ [classification]: dynamic $p.\text{class}$ computation requires:
 - [reference] p^r : $p^r.\text{index}$, $p^r.\text{class}$ (64-bit vector BC^r)
 - [current] p_c : $p_c.\text{index}$
- ➋ [trimming]: find f with different in/out classifications

shader

Rendering & Trimming Algorithms for Dynamic Rules

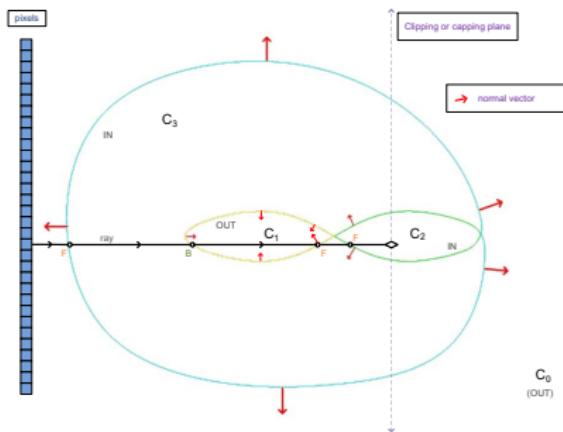
Pipeline

- ➊ [classification]: dynamic $p.\text{class}$ computation requires:
 - [reference] p^r : $p^r.\text{index}$, $p^r.\text{class}$ (64-bit vector BC^r)
 - [current] p_c : $p_c.\text{index}$
- ➋ [trimming]: find f with different in/out classifications

shader

Frame characterization

- [reference]: $\text{frame.class} = 0$
- [current]: $\text{frame.class} = 1$
- [both]: $\text{frame.class} = 2$



Extend Dynamic Trimming for CSG/Capping

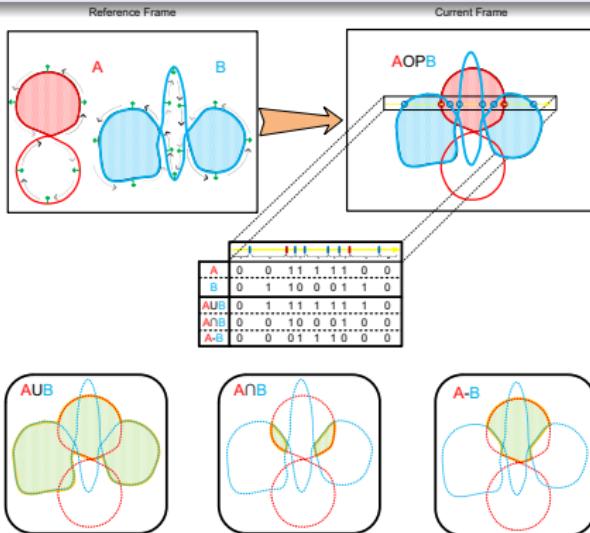
shader

- [solid A]: $\text{frame}.class = 1$ - [solid B]: $\text{frame}.class = 2$
- store/csg operate two classification vectors: $BC_A^r[i] \text{ op } BC_B^r[j]$

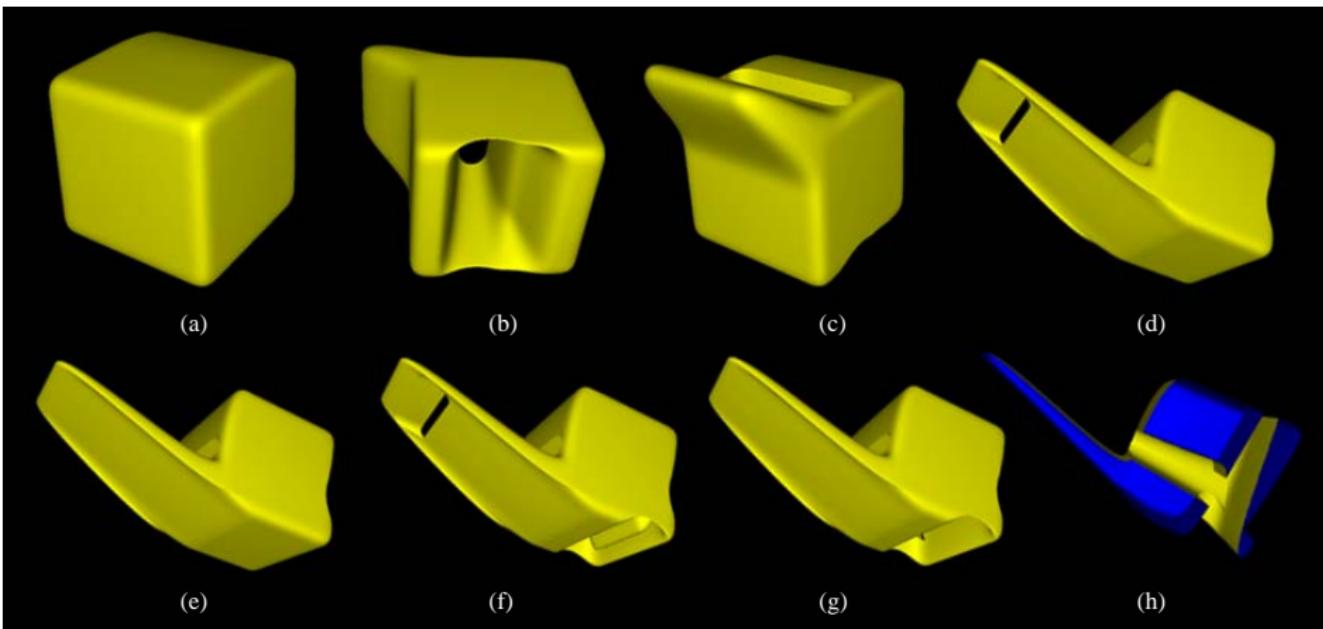
Extend Dynamic Trimming for CSG/Capping

shader

- [solid A]: $\text{frame}.class = 1$ - [solid B]: $\text{frame}.class = 2$
- store/csg operate two classification vectors: $BC_A^r[i] \text{ op } BC_B^r[j]$
- [capping]: subtracting a capping box from our surface



Static & Dynamic Trimming



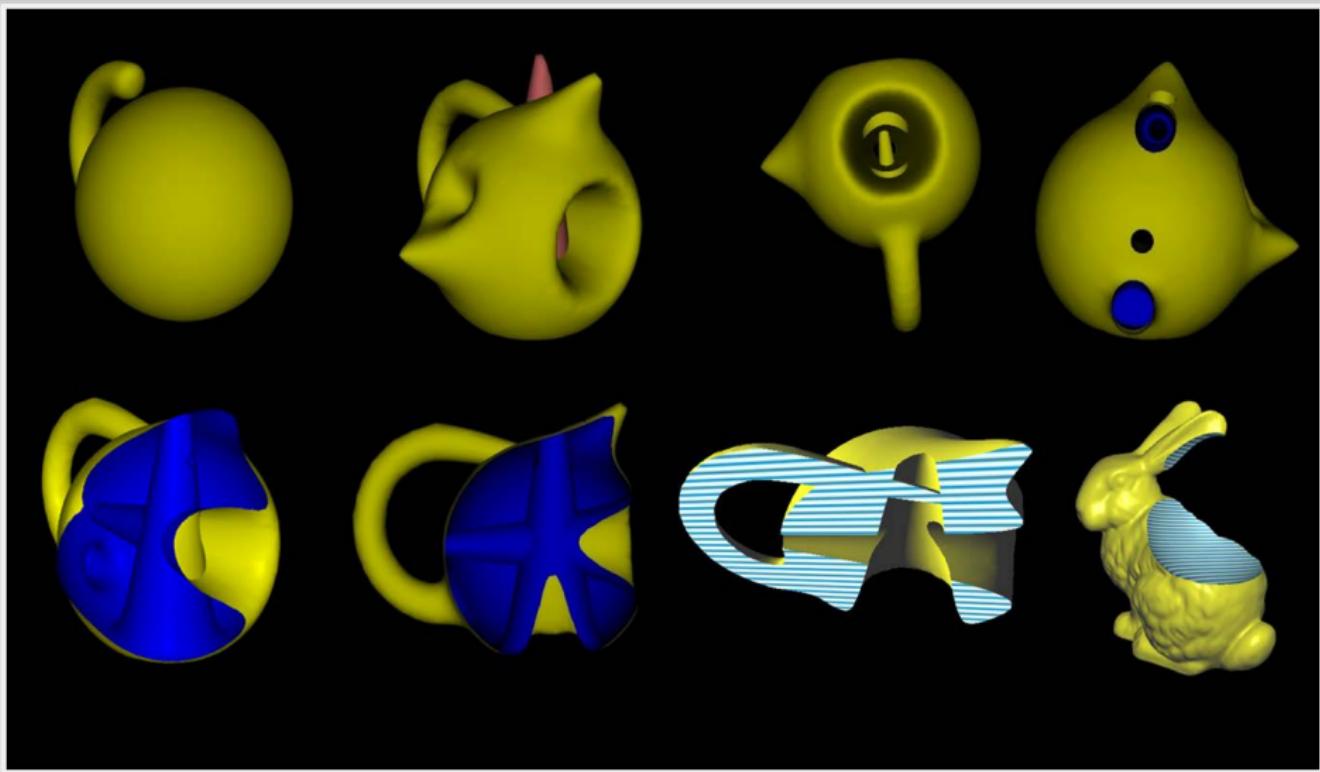
Background
Pose-to-pose segmentation
Pose-to-pose skinning
S-buffer

Eliminating z-fighting flaws
 k^+ -buffer
Rendering self-trimmed solids
Conclusions & Future Work

Problem & Contribution
Winding Number
Static Rules
Dynamic Rules

Rendering Algorithms
Results
Experimental Study
Demonstration

Clipping & Capping



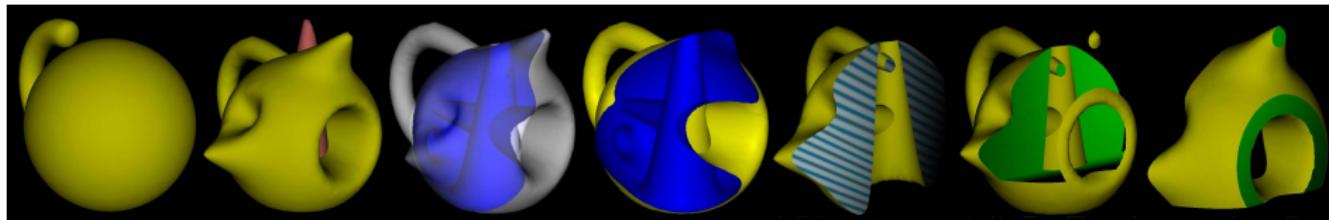
Background
Pose-to-pose segmentation
Pose-to-pose skinning
S-buffer

Eliminating z-fighting flaws
 k^+ -buffer
Rendering self-trimmed solids
Conclusions & Future Work

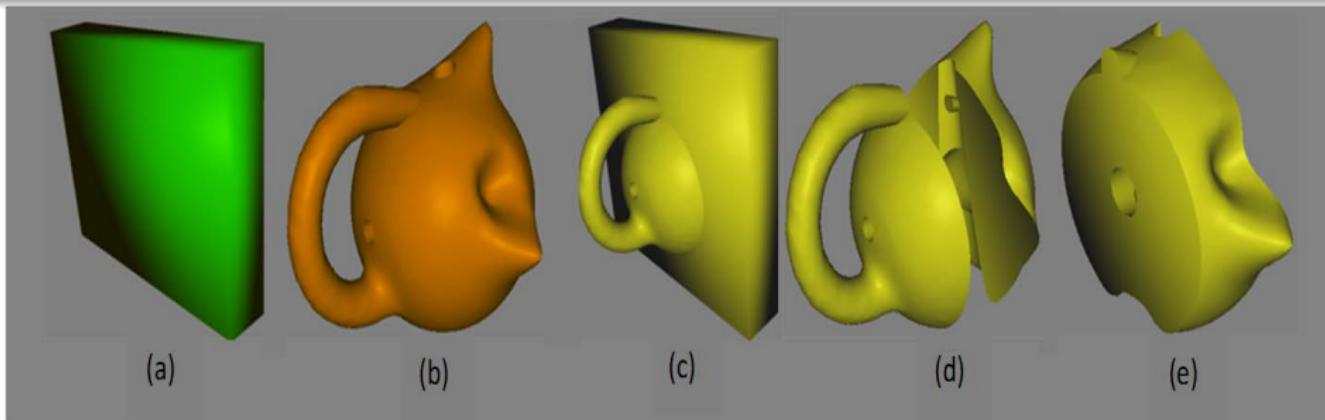
Problem & Contribution
Winding Number
Static Rules
Dynamic Rules

Rendering Algorithms
Results
Experimental Study
Demonstration

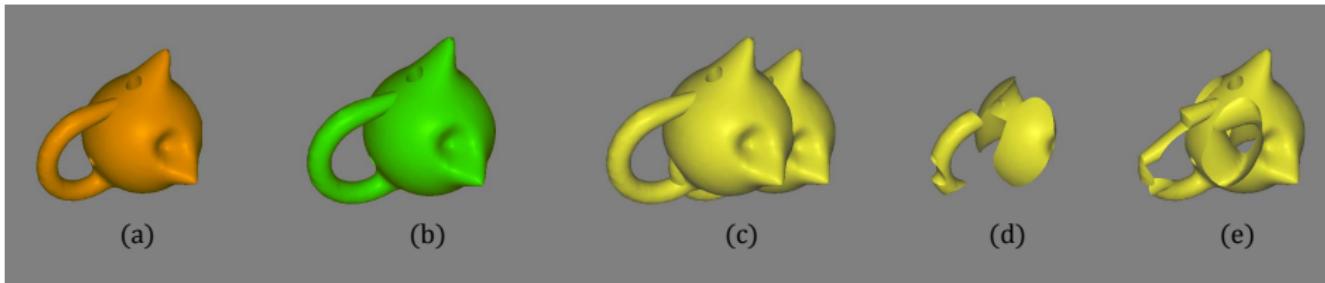
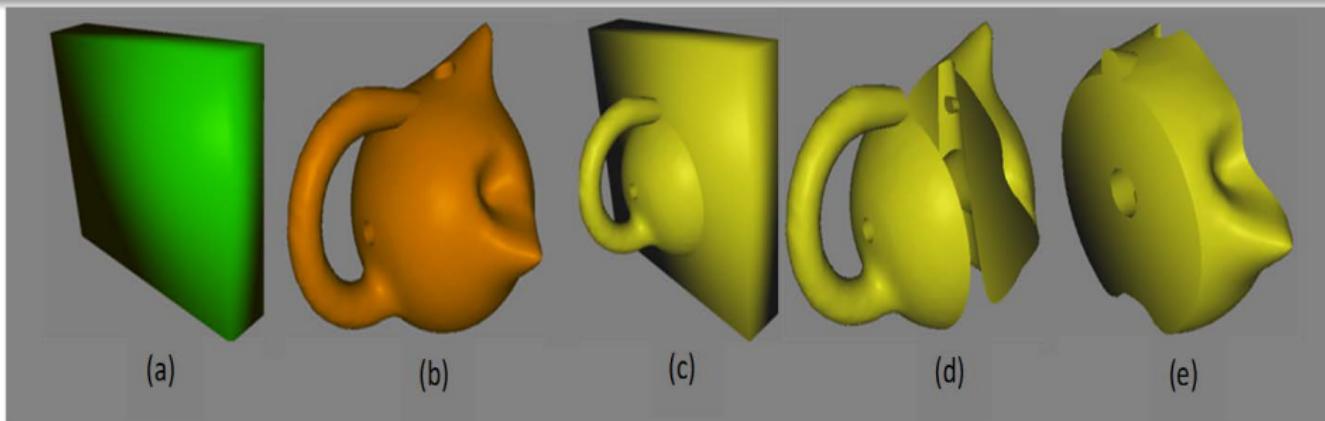
Rendering & Transparency & Clipping & Capping & CSG



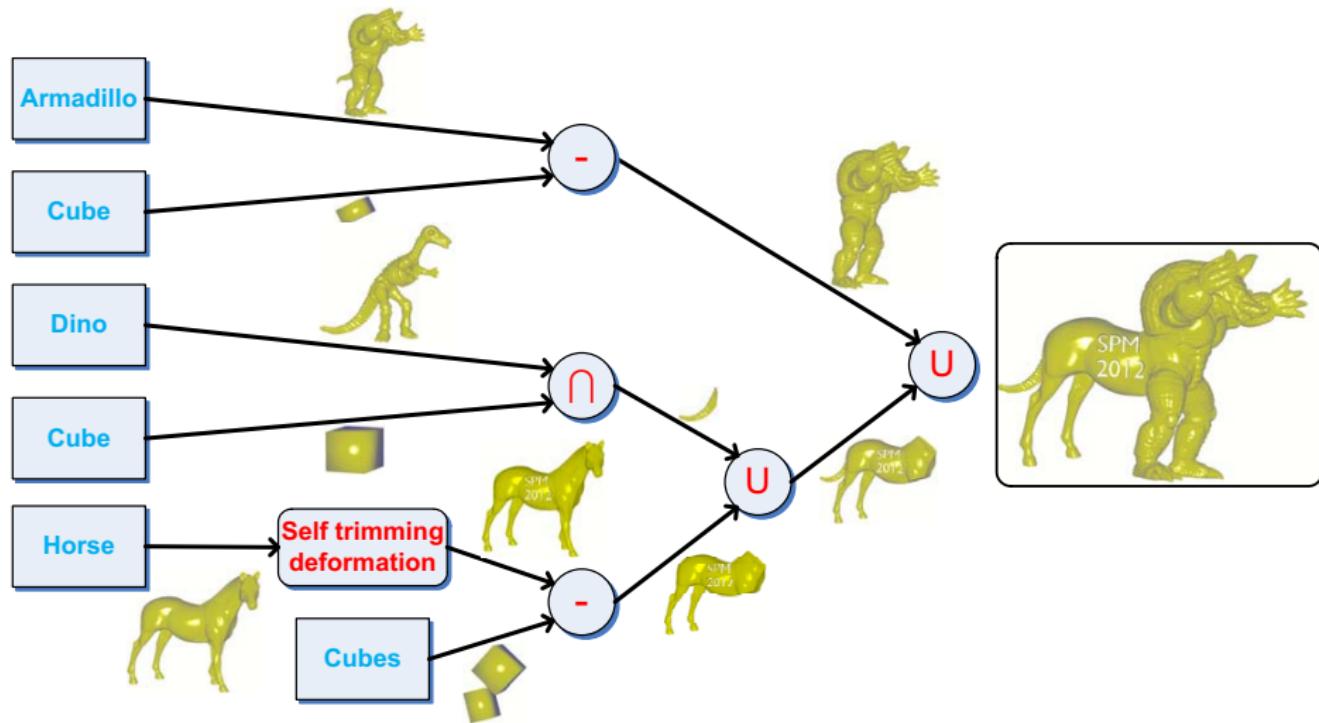
Constructive Solid Geometry



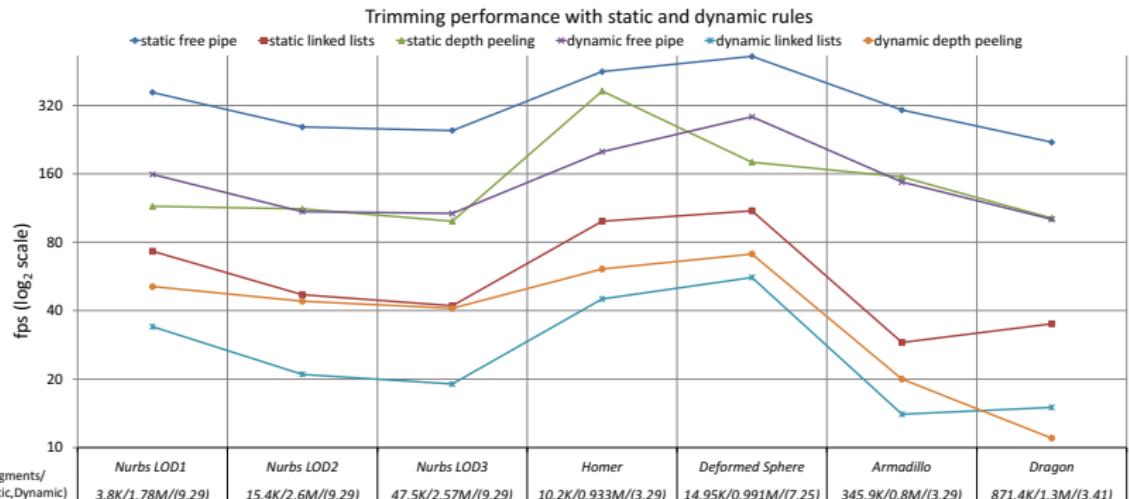
Constructive Solid Geometry



CSG Tree



Performance



Background
Pose-to-pose segmentation
Pose-to-pose skinning
S-buffer

Eliminating z-fighting flaws
 k^+ -buffer
Rendering self-trimmed solids
Conclusions & Future Work

Problem & Contribution
Winding Number
Static Rules
Dynamic Rules

Rendering Algorithms
Results
Experimental Study
Demonstration

Video - Sphere

Background
Pose-to-pose segmentation
Pose-to-pose skinning
S-buffer

Eliminating z-fighting flaws
 k^+ -buffer
Rendering self-trimmed solids
Conclusions & Future Work

Problem & Contribution
Winding Number
Static Rules
Dynamic Rules

Rendering Algorithms
Results
Experimental Study
Demonstration

Video - Nurbs

Outline

- 1 Background and Preliminaries
- 2 Pose partitioning for multi-resolution segmentation of arbitrary MAs
- 3 Pose-to-pose skinning of animated meshes
- 4 S-buffer: Sparsity-aware multi-fragment rendering
- 5 Depth-Fighting Aware Methods for Multi-fragment Rendering
- 6 k^+ -buffer: Fragment synchronized k -buffer
- 7 Direct rendering of self-trimmed surfaces
- 8 Conclusions & Future Work

Conclusions

Segmentation by combining pose partitionings

- handles off-line, real-time and editable MAs
- supports rigid, highly-deformable and hybrid MAs
- enables multi-resolution & variable segmentations
- introduces smooth visualization of segment transitions

Conclusions

Segmentation by combining pose partitionings

- handles off-line, real-time and editable MAs
- supports rigid, highly-deformable and hybrid MAs
- enables multi-resolution & variable segmentations
- introduces smooth visualization of segment transitions

Approximating MAs with pose-to-pose skinning

- handles both linear and non-linear fitting
- supports rigid, highly-deformable and hybrid MAs
- enables pose editing of arbitrary animation frames
- increases skinning quality via novel vertex and weight corrections

Conclusions

Multi-fragment rendering via S-buffer

- two-geometry-pass GPU-accelerated A-buffer that take advantage of
 - the fragment distribution (exact memory allocation)
 - the sparsity of the pixel-space (improved performance)
- organizes storage into variable contiguous regions for each pixel
- integrates into the standard graphics pipeline

Conclusions

Multi-fragment rendering via S-buffer

- two-geometry-pass GPU-accelerated A-buffer that take advantage of
 - the fragment distribution (exact memory allocation)
 - the sparsity of the pixel-space (improved performance)
- organizes storage into variable contiguous regions for each pixel
- integrates into the standard graphics pipeline

Z-fighting free multi-fragment rendering

- introduces robust and approximate algorithms by extending conventional multi-pass rendering methods
- increases multi-pass rendering performance via geometry culling
- integrates into the standard graphics pipeline

Conclusions

Bounded multi-fragment storage using k^+ -buffer

- alleviates prior k -buffer limitations and bottlenecks by exploiting fragment culling and pixel synchronization
- introduces an extension to avoid wasteful memory consumption.
- can also simulate the behavior of Z-buffer or A-buffer

Conclusions

Bounded multi-fragment storage using k^+ -buffer

- alleviates prior k -buffer limitations and bottlenecks by exploiting fragment culling and pixel synchronization
- introduces an extension to avoid wasteful memory consumption.
- can also simulate the behavior of Z-buffer or A-buffer

Real-time (interior) visualization of self-intersecting static/dynamic meshes

- eliminates the cost of computing the self-intersection geometry curves
- introduces static/dynamic rules for the in/out classification
- proposes GPU-based rendering & trimming algorithms
- supports clipping, capping and boolean operations

Directions for future work

Segmentation by combining pose partitionings

- support of time-varying meshes by exploiting vertex mapping
- increased resolution of over-segmentation (cluster similar poses)
- increase the segmentation quality:
 - ① use of sophisticated mesh clustering algorithm
 - ② rectifying the generated segment boundaries
 - ③ retain important per pose partitions (confidence criterion)
- retain visualization coherency between “far-away” poses

Directions for future work

Segmentation by combining pose partitionings

- support of time-varying meshes by exploiting vertex mapping
- increased resolution of over-segmentation (cluster similar poses)
- increase the segmentation quality:
 - ① use of sophisticated mesh clustering algorithm
 - ② rectifying the generated segment boundaries
 - ③ retain important per pose partitions (confidence criterion)
- retain visualization coherency between “far-away” poses

Approximating MAs with pose-to-pose skinning

- reduce the dimensions of the fitting optimization by:
 - ① extracting key-frames
 - ② simplifying the rest-pose
 - ③ performing computations in GPU
- automatically discover the smallest number of suitable joints

Directions for future work

Multi-fragment rendering

- develop a more sophisticated hash function that uniformly divide the non-empty pixels (S-buffer)
- reduce cost of additional accumulation step (S-buffer, k^+ -buffer):
 - ① lower-detailed subdivision of the initial 3D scene
 - ② exploit temporal coherence solutions
- investigate a hybrid depth peeling technique (z-fighting):
 - ① captures a sequence of coplanarity-free layers followed by
 - ② on demand peeling of overlapping fragments
- explore dynamic k^+ -buffer: k value is not the same for all pixels

Directions for future work

Multi-fragment rendering

- develop a more sophisticated hash function that uniformly divide the non-empty pixels (S-buffer)
- reduce cost of additional accumulation step (S-buffer, k^+ -buffer):
 - ① lower-detailed subdivision of the initial 3D scene
 - ② exploit temporal coherence solutions
- investigate a hybrid depth peeling technique (z-fighting):
 - ① captures a sequence of coplanarity-free layers followed by
 - ② on demand peeling of overlapping fragments
- explore dynamic k^+ -buffer: k value is not the same for all pixels

Rendering and trimming of self-crossing solids

- support of fast collision detection
- voxelize the interior and produce actual self-trimmed mesh

My Publications

-  **A. Vasilakis and I. Fudos, Pose Partitioning for Multi-resolution Segmentation of Arbitrary Mesh Animations**, Computer Graphics Forum (Proceedings of Eurographics 2014), vol. 33(2):?-?, 04-2014.
-  **A. Vasilakis and I. Fudos, k^+ -buffer: Fragment Synchronized k -buffer**, Proceedings of the ACM SIGGRAPH Symposium on Interactive 3D Graphics and Games (I3D '14), San Francisco, USA, 03-2014.
-  **A. Vasilakis and I. Fudos, Depth-fighting Aware Methods for Multifragment Rendering**, IEEE Transactions on Visualization and Computer Graphics, vol. 19(6):967-977, 06-2013. [also presented at I3D '13]
-  **J. Rossignac, I. Fudos and A. Vasilakis, Direct Rendering of Boolean Combinations of Self-Trimmed Surfaces**, Computer-Aided Design (Proceedings of Solid and Physical Modeling 2012), Volume 45(2):288-300, 02-2013.
-  **A. Vasilakis and I. Fudos, S-buffer: Sparsity-aware Multi-fragment Rendering**, Proceedings of Eurographics 2012 (Short Papers), 101-104, Cagliari, Italy, 05-2012.
-  **A. Vasilakis and I. Fudos, Z-fighting aware Depth Peeling**, SIGGRAPH 2011 (Posters), Vancouver, Canada, 08-2011.
-  **A. Vasilakis, G. Antonopoulos and I. Fudos, Pose-to-Pose Skinning of Animated Meshes**, Proceedings of Symposium of Computer Animation 2011 (Posters), Vancouver, Canada, 08-2011.
-  **A. Vasilakis and I. Fudos, GPU Rigid Skinning based on a Refined Skeletonization Method**, Computer Animation and Virtual Worlds, vol. 22: 27-46, 01-2011.
-  **A. Vasilakis and I. Fudos, Skeleton-based Rigid Skinning for Character Animation**, Proceedings of the Fourth International Conference on Computer Graphics Theory and Applications (GRAPP '09), 302-308, Lisboa, Portugal, 02-2009.

Acknowledgements

This research has been co-financed by the European Union (European Social Fund ESF) and Greek national funds through the Operational Program Education and Lifelong Learning of the National Strategic Reference Framework (NSRF) - Research Funding Program: Heracleitus II. Investing in knowledge society through the European Social Fund.



European Union
European Social Fund



MINISTRY OF EDUCATION & RELIGIOUS AFFAIRS
MANAGING AUTHORITY

Co-financed by Greece and the European Union



Background
Pose-to-pose segmentation
Pose-to-pose skinning
S-buffer

Eliminating z-fighting flaws
 k^+ -buffer
Rendering self-trimmed solids
Conclusions & Future Work

Conclusions
Future work
My Publications

Acknowledgements
Questions?

Questions?

Thank you!

Z-fighting (more)

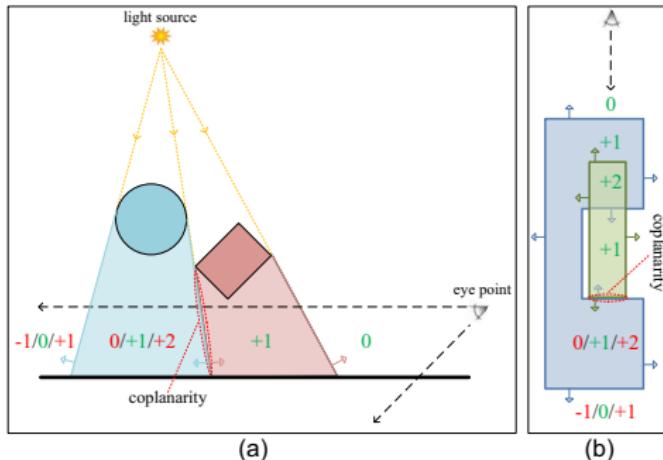
[go back](#)

Figure: Illustrating the values of the popular winding number (used for in/out classifications) when ray casting for (a) shadow volume computations and (b) CSG modeling. Red-painted values highlight erroneous computations (in cases where only one of the two coplanar fragments is successfully captured).

Over-segmentation (more)

Algorithm 1 Over-segmentation(A)

```
1:  $OS(A) \leftarrow \emptyset; VL \leftarrow$  list of all vertices in  $V$ ;  
2:  $k_{max} \leftarrow 0; \forall v \in V : v.segment \leftarrow -1$ ;  
3: Mark all edges as not visited;  
4: while  $VL \neq \emptyset$  do  
5:    $v \leftarrow VL.next$ ;  
6:   if  $v.segment == -1$  then  
7:      $S_{k_{max}} \leftarrow \{v\}; v.segment \leftarrow k_{max}$ ;  
8:      $k \leftarrow k_{max}; k_{max} \leftarrow k_{max} + 1$ ;  
9:   else  
10:     $k \leftarrow v.segment$ ;  
11:   end if  
12:   for each not visited edge  $e_i \leftarrow (v, v_i)$  do  
13:     Mark edge  $e_i$  as visited;  
14:     if  $cav(v) == cav(v_i)$  then  
15:        $S_k \leftarrow S_k \cup \{v_i\}$ ;  
16:        $v_i.segment \leftarrow k$ ;  
17:     end if  
18:   end for  
19:   Remove vertex  $v$  from list  $VL$ ;  
20: end while  
21:  $OS(A) = \{S_0, S_1, \dots, S_{k_{max}-1}\}$ ;
```

Pose-to-pose cleaning (more)

[go back](#)

Algorithm 2 p2p-Cleaning($OS(A), h$)

```

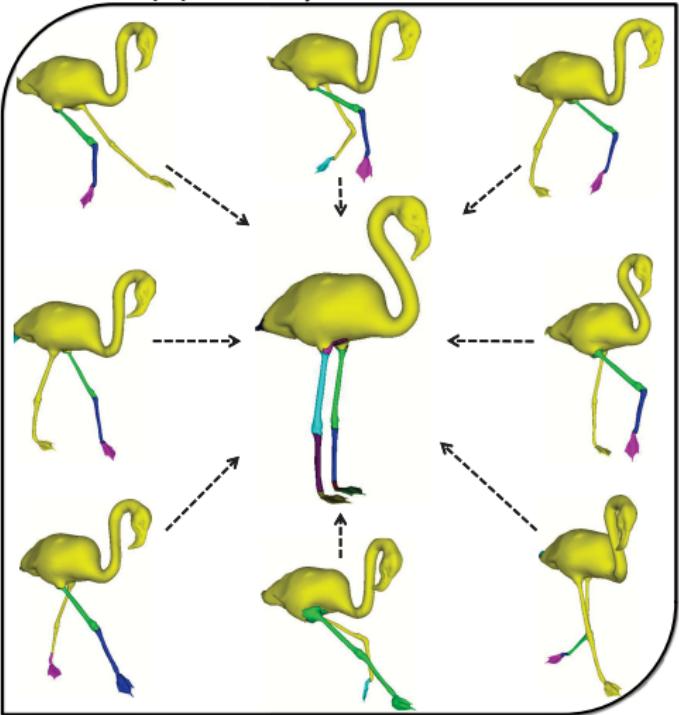
1: for  $p \leftarrow k, 1$  do
2:    $OS(A).computeArea();$ 
3:   for each  $S_A \in OS(A)$  do
4:     if  $Clean(S_A, C^{cav}(S_A)[p], h, p)$  then
5:        $maxArea \leftarrow a(S_A);$ 
6:       for  $S_B \in N(S_A)$  do
7:         if  $maxArea < a(S_B)$  and  $cav(S_A)[1, p - 1] ==$ 
8:            $cav(S_B)[1, p - 1]$  and  $cav(S_A)[p] \neq cav(S_B)[p]$  then
9:              $maxArea \leftarrow a(S_B);$ 
10:            end if
11:          end for
12:          if  $maxArea > a(S_A)$  then
13:             $S_B.Copy(S_A);$  ▷ neighbors,vertices,faces
14:            comment: Leave  $a(S_B)$  unchanged
15:             $OS(A).Remove(S_A);$ 
16:          end if
17:        end for
18:      end for

19: function CLEAN( $S_A, S_P, h, p$ )
20:   if  $a(S_A) > h \cdot a(S_P)$  then
21:     return false;
22:   end if
23:    $S_L = \text{new List<}Vertex\text{>}();$ 
24:   comment: The following is realized with a breadth
25:   first search from  $S_B$  (similar to Algorithm 1)
26:   for each  $v \in segment(S_A, p - 1)$  do
27:      $S_L.Add(v);$ 
28:   end for
29:   return  $a(S_A) \leq h \cdot a(S_L);$ 
30: end function
  
```

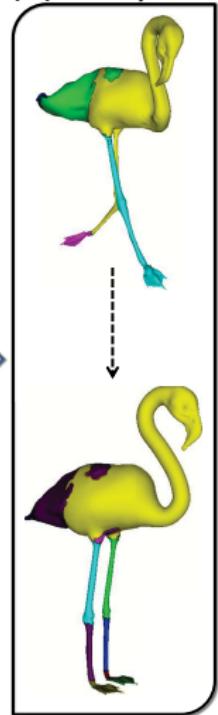
Segmentation (add new poses)

[go back](#)

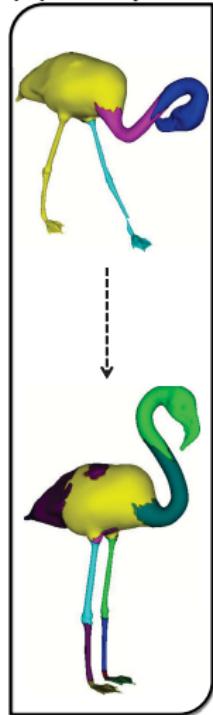
(a) initial pose dataset



(b) add pose

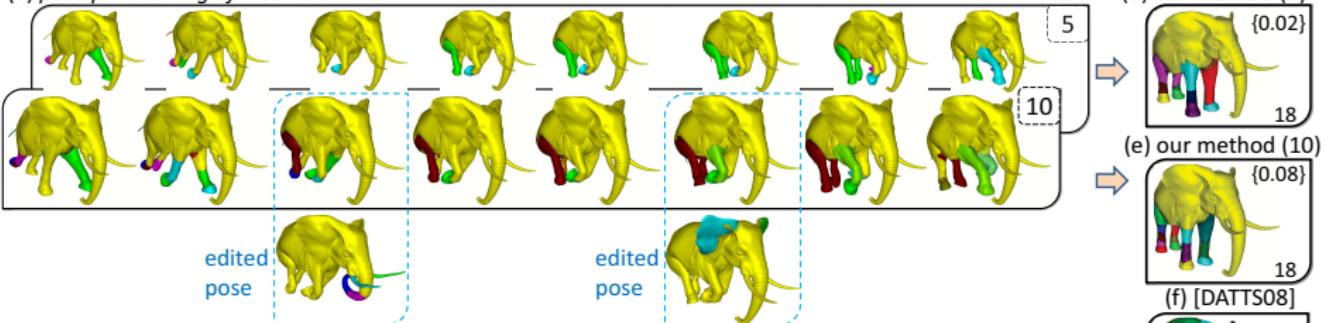


(c) add pose

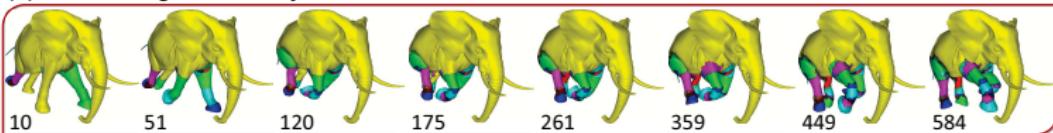


Segmentation (edit poses)

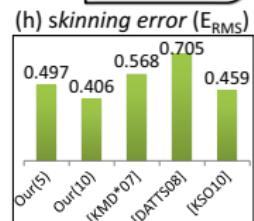
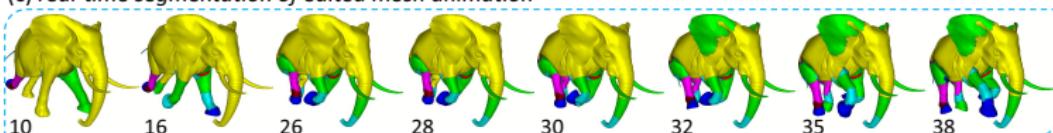
(a) pose partitioning of mesh animation



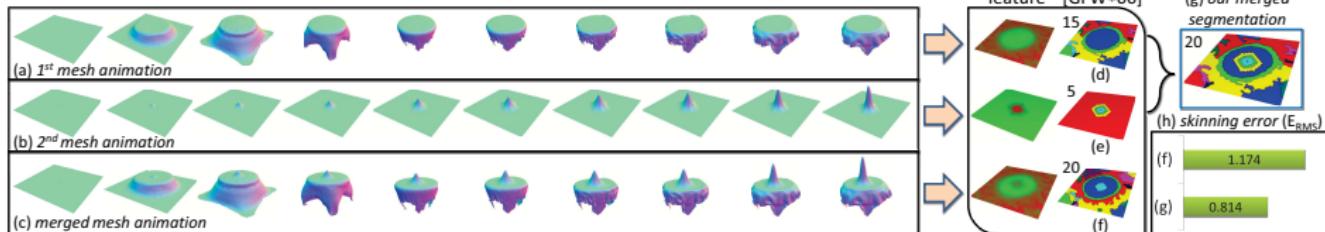
(b) real-time segmentation of mesh animation



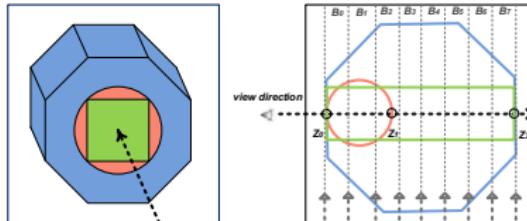
(c) real-time segmentation of edited mesh animation



Segmentation transfer

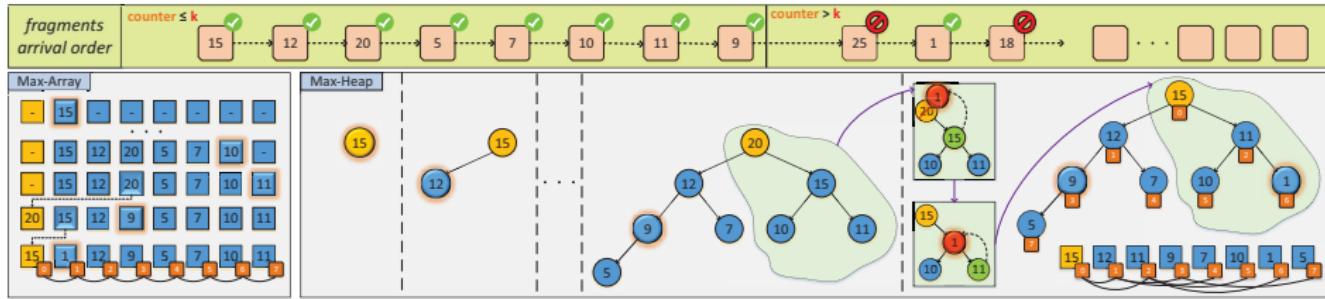


Depth peeling example

[go back](#)


	Iterations			
	1	2	3	4
F2B	[Z]	[Z]	[Z]	
DUAL	[Z]	[Z]	[Z]	
F2B-3P/ F2B-2P	[Z] [Z] [Z]	[Z] [X] [1]	[Z] [X] [1]	[Z] [Z] [2]
DUAL-2P	[Z] [Z] [3] [Z] [Z] [2]	[Z] [X] [1] [Z] [X] [1]		
F2B-LL	[Z → Z → Z]	[Z →]	[Z →]	
DUAL-LL	[Z → Z → Z] [Z → Z →]	[Z →]		
F2B-B	[Z] [Z] [Z] [X] [X] [X]	[Z] [X] [X] [X]	[Z] [Z] [X] [X] [X]	
DUAL-B	[Z] [Z] [Z] [X] [X] [X] [Z] [Z] [X] [X] [X]	[Z] [X] [X] [X]		
AB _{LU} /SB	[Z → Z → Z → Z → Z → Z →]			
KB/KB _{2D} / AB _{3D}	[Z] [Z] [Z] [Z] [Z] [Z] [X] [X] [X]			
KB-Multi		[Z] [Z] [Z] [X] [X] [X] [X] [X]		
BUN	[B ₀] [B ₁] [B ₂] [B ₃] [B ₄] [B ₅] [B ₆] [B ₇]	[Z] [X] [Z] [X] [Z] [X] [Z] [X]		
BUN-LL	[Z → Z → Z] [Z →] [Z → Z → Z →]			

k-buffer store example



k-buffer methods details

Algorithm	Description	Performance	Sorting need		Peeling Accuracy		Memory	
			Rendering Passes	on primitives	on fragments	Max k	32bit Float Precision	Per Pixel Allocation
KB	Initial k -buffer implementation [Callahan2005,Bavoil2007]	1	v	v		8; 16		2k; 4k
KB-Multi	Multi-pass k -buffer [Liu2009a]	1 to k	v	v				2k; 4k
KB-SR	Stencil routed k -buffer [Bavoil2008]	1	v	v	32			3k
KB-PS	k -buffer using pixel synchronization extension [Salvi2013]	1	x	v	-			2k
K'B-Array	k' -buffer using max-array	1	x	v	-			2k + 2
K'B-Heap	k' -buffer using max-heap	1	x	v	-			2k + 2
KB-MDT	Multi depth test scheme [Liu2010,Maule2013]	2	x	x	-			2k
KB-MHA	Memory-hazard-aware k -buffer [Zhang2013]	1	v	v	8; 16		x	2k; 4k
KB-AB _{Array}	k -buffer based on A-buffer (fixed-size arrays)	1	x	v	-			2n + 1
KB-AB _{LL}	k -buffer based on A-buffer (dynamic linked lists) [Yu2012]	1	x	v	-			3f + 1
KB-LL	k -buffer based on linked lists [Yu2012]	1	x	x	-			3f + 6
KB-AB _{SB}	k -buffer based on S-buffer (variable-contiguous regions)	2	x	v	-			2f + 2
K'B-SB	Memory-friendly variation of k' -buffer	2	x	v	-			2f _k + 3
$f(p) = \# \text{fragments at pixel } p[x,y]$		$n = \max_{x,y}\{f(p)\}$				In A ; B, A denotes the layers/memory for the basic method and B for the variation using attribute packing		
$f_k(p) = (f(p) < k) ? f(p) : k$		$f_k(p) \leq k$						

Dynamic Trimming Algorithm

Algorithm 13 TrimRenderDynamic(Pixel p , Fragment f)

```
/* continue until we find the first non trimmed boundary fragment */

1:  $p^c.\text{depth} := C.z;$ 
2: while  $p.\text{color} \neq 0$  do
3:   obtain the next fragment  $f$ :  $f.z > p^c.\text{depth}$ 
4:    $p^c.\text{depth} := f.z;$ 
5:    $w_f := (f \text{ is front facing}) ? 1 : -1$ 
6:   if  $frame.\text{class} \neq 1$  then
7:      $p^r.\text{class} := BC^r[+ + p^r.\text{count}]$ ;
8:      $p^r.\text{index} := p^r.\text{index} + w_f$ ;
9:   end if
10:  if  $frame.\text{class} > 0$  then
11:     $p^c.\text{index} := p^c.\text{index} + w_f$ ;
12:     $CL_b := cCharP$ ;
13:     $CL_a := dynamic\_rule(p^c.\text{index}, p^r.\text{index}, p^c.\text{class})$ ;
14:     $p^c.\text{class} := CL_a$ ;
15:    if  $CL_a \neq CL_b$  then
16:       $p.\text{color} := f.\text{color}$ ;
17:    end if
18:  end if
```

CSG Rendering Algorithm

[go back](#)

Algorithm 14 RenderCSG(Pixel p , Fragment f)

/* continue until we find the first non trimmed boundary fragment */

```
1:  $p^c.\text{depth} := C.z;$ 
2: while  $p.\text{color} \neq 0$  do
3:   obtain the next fragment  $f$ :  $f.z > p^c.\text{depth}$ 
4:    $p^c.\text{depth} := f.z;$ 
5:   if  $frame.\text{class} == 1$  then
6:      $p_B^r.\text{class} := BC_B^r[++Count_B^r];$ 
7:   else if  $frame.\text{class} == 2$  then
8:      $p_A^r.\text{class} := BC_A^r[++Count_A^r];$ 
9:   end if
10:   $CL_b := BC_{AopB}^r[Count_{AopB}^c++];$ 
11:   $CL_a := p_A^r.\text{class } LOGIC\_OP p_B^r.\text{class};$ 
12:   $BC_{AopB}^r[Count_{AopB}^c] := CL_a;$ 
13:  if  $CL_a \neq CL_b$  then
14:     $p.\text{color} := f.\text{color};$ 
15:  end if
16: end while
```
